Package 'DPQ'

September 24, 2019

Title Density, Probability, Quantile ('DPQ') Computations

Version 0.3-3

Date 2019-09-24

Description Computations for approximations and alternatives for the 'DPQ'
(Density (pdf), Probability (cdf) and Quantile) functions for probability
distributions in R.
Primary focus is on (central and non-central) beta, gamma and related
distributions such as the chi-squared, F, and t.
--
This is for the use of researchers in these numerical approximation
implementations, notably for my own use in order to improve R`s own
pbeta(), qgamma(), ..., etc: {``dpq'''-functions}.
-- We plan to complement with 'DPQmpfr' to be suggested later.

Depends R (>= 3.5.0)

Imports stats, graphics, methods, utils, sfsmisc

Suggests Rmpfr, Matrix, mgcv, scatterplot3d, akima

SuggestsNote Matrix only for its ``test-tools-1.R''; mgcv,scatt..,akima:
some tests/

License GPL (>= 2)

Encoding UTF-8

Author Martin Maechler [aut, cre] (<https://orcid.org/0000-0002-8685-9910>),
Morten Welinder [ctb] (pgamma C code, see PR#7307, Jan. 2005),
Wolfgang Viechtbauer [ctb] (dtWV(), 2002),
Ross Ihaka [ctb] (src/qchisq_appr.c),
R-core [ctb] (src/{dpq.h, algdiv.c, pnchisq.c})

Maintainer Martin Maechler <maechler@stat.math.ethz.ch>

Repository R-Forge

Repository/R-Forge/Project specfun

Repository/R-Forge/Revision 111

Repository/R-Forge/DateTimeStamp 2019-09-24 13:38:29

Date/Publication 2019-09-24 13:38:29

R topics documented:

DPQ-package   *Density, Probability, Quantile ('DPQ') Computations*

Description

Computations for approximations and alternatives for the 'DPQ' (Density (pdf), Probability (cdf) and Quantile) functions for probability distributions in R. Primary focus is on (central and non-central) beta, gamma and related distributions such as the chi-squared, F, and t. – This is for the use of researchers in these numerical approximation implementations, notably for my own use in order to improve R's own pbeta(), qgamma(), ..., etc: '"dpq"'-functions. – We plan to complement with 'DPQmpfr' to be suggested later.

Details

The DESCRIPTION file:

Package:	DPQ
Title:	Density, Probability, Quantile ('DPQ') Computations
Version:	0.3-3
Date:	2019-09-24
Authors@R:	c(person("Martin","Maechler", role=c("aut","cre"), email="maechler@stat.math.e
Description:	Computations for approximations and alternatives for the 'DPQ' (Density (pdf), F
Depends:	R (>= 3.5.0)
Imports:	stats, graphics, methods, utils, sfsmisc
Suggests:	Rmpfr, Matrix, mgcv, scatterplot3d, akima
SuggestsNote:	Matrix only for its "test-tools-1.R"; mgcv,scatt..,akima: some tests/
License:	GPL (>= 2)
Encoding:	UTF-8
Author:	Martin Maechler [aut, cre] (<https://orcid.org/0000-0002-8685-9910>), Morten W
Maintainer:	Martin Maechler <maechler@stat.math.ethz.ch>
Repository:	R-Forge
Repository/R-Forge/Project:	specfun
Repository/R-Forge/Revision:	111
Repository/R-Forge/DateTimeStamp:	2019-09-24 13:38:29
Date/Publication:	2019-09-24 13:38:29

Index of help topics:

DPQ-package	Density, Probability, Quantile ('DPQ') Computations
M_LN2	Numerical Utilities - Functions, Constants
algdiv	Compute log(gamma(b)/gamma(a+b)) when b >= 8
b_chi	Compute E[chi_nu]/sqrt(nu) useful for t- and chi-Distributions
bd0	Utility Functions for 'dgamma()' - Pure R Versions
dgamma.R	Gamma Density Function Alternatives
dnchisqR	Approximations of the (Noncentral) Chisquared Density
dntJKBf1	Non-central t-Distribution Density - Algorithms and Approximations
dtWV	Noncentral t Distribution Density by W.V.
format01prec	Format Numbers in [0,1] with "Precise" Result
lbetaM	(Log) Beta Approximations
lgamma1p	Accurate 'log(gamma(a+1))'
log1mexp	Compute f(a) = log(1 - exp(-a)) Numerically Optimally
log1pmx	Accurate 'log(1+x) - x'
logcf	Continued Fraction Approximation of Log-Related Series
logspace.add	Logspace Arithmetix - Addition and Subtraction
lssum	Compute Logarithm of a Sum with Signed Large Summands
lsum	Properly Compute the Logarithm of a Sum

newton	Simple R level Newton Algorithm, Mostly for Didactical Reasons
pbetaRv1	Pure R Implementation of Old pbeta()
pchisqV	Wienergerm Approximations to (Non-Central) Chisquare Probabilities
pl2curves	Plot 2 Noncentral Distribution Curves for Visual Comparison
pnbetaAppr2	Noncentral Beta Probabilities
pnchi1sq	(Probabilities of Non-Central Chisquared Distribution for Special Cases
pnchisq	(Approximate) Probabilities of Non-Central Chisquared Distribution
pntR	Non-central t Probability Distribution - Algorithms and Approximations
ppoisErr	Direct Computation of 'ppois()' Poisson Distribution Probabilities
qbetaAppr	Compute (Approximate) Quantiles of the Beta Distribution
qchisqAppr	Compute Approximate Quantiles of the Chi-Squared Distribution
qgammaAppr	Compute (Approximate) Quantiles of the Gamma Distribution
qnchisqAppr	Compute Approximate Quantiles of Noncentral Chi-Squared Distribution
qnormAppr	Approximations to 'qnorm()', i.e., z_alpha
qtAppr	Compute Approximate Quantiles of Non-Central t Distribution
r_pois	Compute Relative Size of i-th term of Poisson Distribution Series

Further information is available in the following vignettes:

| Noncentral-Chisq | Noncentral Chi-Squared Probabilities – Algorithms in R (source) |
| comp-beta | Computing Beta(a,b) for Large Arguments (source) |

Author(s)

Principal author and maintainer: Martin Maechler <maechler@stat.math.ethz.ch>

See Also

The package **DPQmpfr** (not yet on CRAN), which builds on this package and on **Rmpfr**.

Examples

```
"FIXME"
```

---

algdiv *Compute log(gamma(b)/gamma(a+b)) when b >= 8*

Description

Computes

$$algdiv(a,b) := \log \frac{\Gamma(b)}{\Gamma(a+b)} = \log \Gamma(b) - \log \Gamma(a+b) = \texttt{lgamma(b) - lgamma(a+b)}$$

in a numerically stable way.

This is an auxiliary function ins TOMS 708 implementation of `pbeta`, aka the incomplete beta function ratio.

Usage

```
algdiv(a, b)
```

Arguments

a, b   numeric vectors which will be recycled to the same length.

Details

Note that this is also useful to compute the Beta function

$$B(a,b) = \frac{\Gamma(a)\Gamma(b)}{\Gamma(a+b)}.$$

Clearly,
$$\log B(a,b) = \log \Gamma(a) + algdiv(a,b) = \log \Gamma(a) - logQab(a,b)$$

```
 In our ../tests/qbeta-dist.R we look into computing  log(p*Beta(p,q)) accurately for p << q
        ---------------------
 We are proposing a nice solution there.

 How is this related to algdiv() ?
```

Value

a numeric vector of length `max(length(a),length(b))` (if neither is of length 0, in which case the result has length 0 as well).

Author(s)

Didonato, A. and Morris, A., Jr, (1992); `algdiv()`'s C version from the R sources, authored by the R core team; C and R interface: Martin Maechler

References

Didonato, A. and Morris, A., Jr, (1992) Algorithm 708: Significant digit computation of the incomplete beta function ratios, *ACM Transactions on Mathematical Software* **18**, 360–373.

See Also

gamma, beta; my own logQab_asy().

Examples

```
Qab <- algdiv(2:3, 8:14)
cbind(a = 2:3, b = 8:14, Qab) # recycling with a warning

## algdiv()  and my  logQab_asy()  give *very* similar results for largish b:
all.equal( -  algdiv(3, 100),
            logQab_asy(3, 100), tol=0) # 1.283e-16 !!
(lQab <- logQab_asy(3, 1e10))
## relative error
1 + lQab/ algdiv(3, 1e10) # 0 (64b F 30 Linux; 2019-08-15)
```

---

b_chi   *Compute $E[\chi\_\nu]/\sqrt{\nu}$ useful for t- and chi-Distributions*

Description

$$b_\chi(\nu) := E[\chi(\nu)]/\sqrt{\nu} = \frac{\sqrt{2/\nu}\Gamma((\nu+1)/2)}{\Gamma(\nu/2)},$$

where $\chi(\nu)$ denotes a chi-distributed random variable, i.e., the square of a chi-squared variable, and $\Gamma(z)$ is the Gamma function, gamma() in R.

This is a relatively important auxiliary function when computing with non-central t distribution functions and approximations, specifically see Johnson et al.(1994), p.520, after (31.26a).

Its logarithm,

$$lb_\chi(\nu) := log\big(\frac{\sqrt{2/\nu}\Gamma((\nu+1)/2)}{\Gamma(\nu/2)}\big),$$

is even easier to compute via lgamma and log, and I have used Maple to derive an asymptotic expansion in $\frac{1}{\nu}$ as well.

Note that $lb_\chi(\nu)$ also appears in the formula for the t-density (dt) and distribution (tail) functions.

Usage

```
b_chi     (nu, one.minus = FALSE, c1 = 341, c2 = 1000)
b_chiAsymp(nu, order = 2, one.minus = FALSE)
#lb_chi    (nu, ......) # not yet
lb_chiAsymp(nu, order)

c_dt(nu)     # warning("FIXME: current c_dt() is poor -- base it on lb_chi(nu) !")
c_dtAsymp(nu)  # deprecated in favour of lb_chi(nu)
c_pt(nu)     # warning("use better c_dt()") %---> FIXME deprecate even stronger ?
```

Arguments

nu	non-negative numeric vector of degrees of freedom.
one.minus	logical indicating if $1 - b()$ should be returned instead of $b()$.
c1, c2	boundaries for different approximation intervals used: for 0 < nu <= c1, internal b1() is used, for c1 < nu <= c2, internal b2() is used, and for c2 < nu, the b_chiAsymp() function is used, (and you can use that explicitly, also for smaller nu). FIXME: c1 and c2 were defined when the only asymptotic expansion known to me was the order = 2 one. A future version of b_chi will *very likely* use b_chiAsymp(*,order) for higher orders, and the c1 and c2 arguments will change, possibly be abolished.
order	the polynomial order in $\frac{1}{\nu}$ of the asymptotic expansion of $b_\chi(\nu)$ for $\nu \to \infty$. The default, order = 2 corresponds to the order you can get out of the Abramowitz and Stegun (6.1.47) formula. Higher order expansions were derived using Maple by Martin Maechler in 2002, see below, but implemented in b_chiAsymp() only in 2018.

Details

One can see that b_chi() has the properties of a CDF of a continuous positive random variable: It grows monotonely from $b_\chi(0) = 0$ to (asymptotically) one. Specifically, for large nu, b_chi(nu) = b_chiAsymp(nu) and

$$1 - b_\chi(\nu) \sim \frac{1}{4\nu}.$$

More accurately, derived from Abramowitz and Stegun, 6.1.47 (p.257) for a= 1/2, b=0,

$$\Gamma(z + 1/2)/\Gamma(z) \sim \sqrt(z) * (1 - 1/(8z) + 1/(128z^2) + O(1/z^3)),$$

and applied for $b_\chi(\nu)$ with $z = \nu/2$, we get

$$b_\chi(\nu) \sim 1 - (1/(4\nu) * (1 - 1/(8\nu)) + O(\nu^{-3})),$$

which has been implemented in b_chiAsymp(*,order=2) in 1999.

Even more accurately, Martin Maechler, used Maple to derive an asymptotic expansion up to order 15, here reported up to order 5, namely with $r := \frac{1}{4\nu}$,

$$b_\chi(\nu) = c_\chi(r) = 1 - r + \frac{1}{2}r^2 + \frac{5}{2}r^3 - \frac{21}{8}r^4 - \frac{399}{8}r^5 + O(r^6).$$

Value

a numeric vector of the same length as nu.

Author(s)

Martin Maechler

References

Johnson, Kotz, Balakrishnan (1995) *Continuous Univariate Distributions*, Vol 2, 2nd Edition; Wiley.
Formula on page 520, after (31.26a)

Abramowitz, M. and Stegun, I. A. (1972) *Handbook of Mathematical Functions*. New York: Dover.
https://en.wikipedia.org/wiki/Abramowitz_and_Stegun provides links to the full text which is in public domain.

See Also

The t-distribution (base R) page pt.

Examples

```
curve(b_chi, 0, 20); abline(h=0:1, v=0, lty=3)
r <- curve(b_chi, 1e-10, 1e5, log="x")
with(r, lines(x, b_chi(x, one.minus=TRUE), col = 2))

## Zoom in to c1-region
rc1 <- curve(b_chi, 340.5, 341.5, n=1001)# nothing to see
e <- 1e-3; curve(b_chi, 341-e, 341+e, n=1001) # nothing
e <- 1e-5; curve(b_chi, 341-e, 341+e, n=1001) # see noise, but no jump
e <- 1e-7; curve(b_chi, 341-e, 341+e, n=1001) # see float "granularity"+"jump"

## Zoom in to c2-region
rc2 <- curve(b_chi, 999.5, 1001.5, n=1001) # nothing visible
e <- 1e-3; curve(b_chi, 1000-e, 1000+e, n=1001) # clear small jump
c2 <- 1500
e <- 1e-3; curve(b_chi(x,c2=c2), c2-e, c2+e, n=1001)# still
## - - - -
c2 <- 3000
e <- 1e-3; curve(b_chi(x,c2=c2), c2-e, c2+e, n=1001)# ok asymp clearly better!!
curve(b_chiAsymp, add=TRUE, col=adjustcolor("red", 1/3), lwd=3)
if(requireNamespace("Rmpfr")) {
 xm <- Rmpfr::seqMpfr(c2-e, c2+e, length.out=1000)

}
## - - - -
c2 <- 4000
e <- 1e-3; curve(b_chi(x,c2=c2), c2-e, c2+e, n=1001)# ok asymp clearly better!!
curve(b_chiAsymp, add=TRUE, col=adjustcolor("red", 1/3), lwd=3)

grCol <- adjustcolor("forest green", 1/2)
curve(b_chi,                    1/2, 1e11, log="x")
curve(b_chiAsymp, add = TRUE, col = grCol, lwd = 3)
## 1-b(nu) ~= 1/(4 nu) a power function <==> linear in log-log scale:
curve(b_chi(x, one.minus=TRUE), 1/2, 1e11, log="xy")
curve(b_chiAsymp(x, one.minus=TRUE), add = TRUE, col = grCol, lwd = 3)
```

---

dchisqApprox	*Approximations of the (Noncentral) Chisquared Density*

Description

Compute the density function $f(x, *)$ of the (noncentral) chisquared distribution.

Usage

```
dnchisqR      (x, df, ncp, log = FALSE,
               eps = 5e-15, termSml = 1e-10, ncpLarge = 1000)
dnchisqBessel(x, df, ncp, log = FALSE)
dchisqAsym    (x, df, ncp, log = FALSE)
dnoncentchisq(x, df, ncp, kmax = floor(ncp/2 + 5 * (ncp/2)^0.5))
```

Arguments

x	non-negative numeric vector.
df	degrees of freedom (parameter), a positive number.
ncp	non-centrality parameter δ; ....
log	logical indicating if the result is desired on the log scale.
eps	positive convergence tolerance for the series expansion: Terms are added while term * q > (1-q)*eps, where q is the term's multiplication factor.
termSml	positive tolerance: in the series expansion, terms are added to the sum as long as they are not smaller than termSml * sum even when convergence according to eps had occured. This was not part of the original C code, but was added later for safeguarding against infinite loops, from PR#14105, e.g., for dchisq(2000,2,1000).
ncpLarge	in the case where mid underflows to 0, when log is true, orncp >= ncpLarge, use a central approximation. In theory, an optimal choice of ncpLarge would not be arbitrarily set at 1000 (hardwired in R's dchisq() here), but possibly also depend on x or df.
kmax	the number of terms in the sum for dnoncentchisq().

Details

dnchisqR() is a pure R implementation of R's own C implementation in the sources, 'R/src/nmath/dnchisq.c', additionally exposing the three "tuning parameters" eps, termSml, and ncpLarge.

dnchisqBessel() implements Fisher(1928)'s exact closed form formula based on the Bessel function I_{nu}, i.e., R's besselI() function; specifically formula (29.4) in Johnson et al. (1995).

dchisqAsym() is the simple asymptotic approximation from Abramowitz and Stegun's formula 26.4.27, p. 942.

dnoncentchisq() uses the (typically defining) infinite series expansion directly, with truncation at kmax, and terms t_k which are products of a Poisson probability and a central chisquare density, i.e., terms t.k := dpois(k,lambda = ncp/2) * dchisq(x,df = 2*k + df) for k = 0,1,...,kmax.

Value

numeric vector similar to x, containing the (logged if log=TRUE) values of the density $f(x, *)$.

Note

These functions are mostly of historical interest, notably as R's dchisq() was not always very accurate in the noncentral case, i.e., for ncp > 0.

Note

R's dchisq() is typically more uniformly accurate than the approximations nowadays, apart from dnchisqR() which should behave the same. There may occasionally exist small differences between dnchisqR(x,*) and dchisq(x,*) for the same parameters.

Author(s)

Martin Maechler, April 2008

References

Abramowitz, M. and Stegun, I. A. (1972) *Handbook of Mathematical Functions*. New York: Dover. https://en.wikipedia.org/wiki/Abramowitz_and_Stegun provides links to the full text which is in public domain.

Johnson, N.L., Kotz, S. and Balakrishnan, N. (1995) Continuous Univariate Distributions Vol~2, 2nd ed.; Wiley.
Chapter 29, Section *3 Distribution*, (29.4), p. 436.

See Also

R's own dchisq().

Examples

```
x <- sort(outer(c(1,2,5), 2^(-4:5)))
fRR <- dchisq  (x, 10, 2)
f.R <- dnchisqR(x, 10, 2)
all.equal(fRR, f.R, tol = 0) # 64bit Lnx (F 30): 1.723897e-16
stopifnot(all.equal(fRR, f.R, tol = 4e-15))
```

---

dgamma-utils   *Utility Functions for* dgamma() – *Pure R Versions*

Description

R transcriptions of the C code utility functions for dgamma() and similar "base" density functions by Catherine Loader.

Usage

```
dpois_raw(x, lambda, log)
bd0(x, np, verbose = getOption("verbose"))
stirlerr(n)
```

Arguments

x, n	a number ([numeric](#)).
lambda, np	a number ([numeric](#); distrubution parameters.
log	logical indicating if the log-density should be returned, otherwise the density at x.
verbose	logical indicating if some information about the computations are to be printed.

Value

a number

Author(s)

Martin Maechler

See Also

[dgamma](#), [dpois](#).

Examples

```
n <- seq(1, 50, by=1/4)
plot(n, sapply(n, stirlerr), type = "b", log="xy")
```

---

dgamma.R	*Gamma Density Function Alternatives*

Description

dgamma.R() is aimed to be an R level "clone" of R's C level implementation [dgamma](#) (from package **stats**).

Usage

```
dgamma.R(x, shape, scale = 1, log)
```

Arguments

x	non-negative numeric vector.
shape	non-negative shape parameter of the Gamma distribution.
scale	positive scale parameter; note we do not see the need to have a rate parameter as the standard R function.
log	logical indicating if the result is desired on the log scale.

Value

numeric vector of the same length as x (which may have to be thought of recycled along shape and/or scale.

Author(s)

Martin Maechler

See Also

(As R's C code) this depends crucially on the "workhorse" function dpois_raw().

Examples

```
## TODO: ... regular case .. use all.equal() ...


## PR#17577 - dgamma(x, shape)  for shape < 1 (=> +Inf at x=0) and very small x
stopifnot(exprs = {
    all.equal(dgamma.R(2^-1027, shape = .99 , log=TRUE), 7.1127667376, tol=1e-10)
    all.equal(dgamma.R(2^-1031, shape = 1e-2, log=TRUE), 702.8889158,  tol=1e-10)
    all.equal(dgamma.R(2^-1048, shape = 1e-7, log=TRUE), 710.30007699, tol=1e-10)
    all.equal(dgamma.R(2^-1048, shape = 1e-7, scale = 1e-315, log=TRUE),
              709.96858768, tol=1e-10)
})
## R's dgamma() gave all Inf in R <= 3.6.1 [and still there in 32-bit Windows !]
```

---

dnt   *Non-central t-Distribution Density - Algorithms and Approximations*

Description

dntJKBf1 implements the summation formulas of Johnson, Kotz and Balakrishnan (1995), (31.15) on page 516 and (31.15') on p.519, the latter being typo-corrected for a missing factor $1/j!$.

dntJKBf() is Vectorize(dntJKBf1,c("x","df","ncp")), i.e., works vectorized in all three main arguments x, df and ncp.

The functions .dntJKBch1() and .dntJKBch() are only there for didactical reasons allowing to check that indeed formula (31.15') in the reference is missing a $j!$ factor in the denominator.

The dntJKBf*() functions are written to also work with arbitrary precise numbers of class "mpfr" (from package **Rmpfr**) as arguments.

Usage

```
dntJKBf1(x, df, ncp, log = FALSE, M = 1000)
dntJKBf (x, df, ncp, log = FALSE, M = 1000)

## The "checking" versions, only for proving correctness of formula:
.dntJKBch1(x, df, ncp, log = FALSE, M = 1000, check=FALSE, tol.check = 1e-7)
.dntJKBch (x, df, ncp, log = FALSE, M = 1000, check=FALSE, tol.check = 1e-7)
```

Arguments

x, df, ncp	see R's dt(); note that each can be of class "mpfr".
log	as in dt(), a logical indicating if $\log(f(x, *))$ should be returned instead of $f(x, *)$.
M	the number of terms to be used, a positive integer.
check	logical indicating if checks of the formula equialities should be done.
tol.check	tolerance to be used for all.equal() when check is true.

Details

How to choose M optimally has not been investigated yet.

Note that relatedly,

R's source code 'R/src/nmath/dnt.c' has claimed from 2003 till 2014 but **wrongly** that the non-central t density $f(x, *)$ is

```
   f(x, df, ncp) =
 df^(df/2) * exp(-.5*ncp^2) /
 (sqrt(pi)*gamma(df/2)*(df+x^2)^((df+1)/2)) *
   sum_{k=0}^Inf gamma((df + k + df)/2)*ncp^k / prod(1:k)*(2*x^2/(df+x^2))^(k/2) .
```

These functions (and this help page prove that it was wrong.

Value

a number for dntJKBf1() and .dntJKBch1().

a numeric vector of the same length as the maximum of the lengths of x,df,ncp for dntJKBf() and .dntJKBch().

Author(s)

Martin Maechler

References

Johnson, N.L., Kotz, S. and Balakrishnan, N. (1995) Continuous Univariate Distributions Vol~2, 2nd ed.; Wiley.
Chapter 31, Section *5 Distribution Function*, p.514 ff

See Also

dt.

Examples

```
tt <- seq(0, 10, len = 21)
ncp <- seq(0, 6, len = 31)
dt3R   <- outer(tt, ncp, dt,      df = 3)
dt3JKB <- outer(tt, ncp, dntJKBf, df = 3)
all.equal(dt3R, dt3JKB)


x <- seq(-1,12, by=1/16)
```

```
fx  <- dt(x, df=3, ncp=5)
re1 <- 1 - .dntJKBch(x, df=3, ncp=5) / fx # with warnings
re2 <- 1 -  dntJKBf (x, df=3, ncp=5) / fx
stopifnot(all.equal(re1[!is.na(re1)], re2[!is.na(re1)], tol=1e-6))
matplot(x, log10(abs(cbind(re1, re2))), type = "o", cex = 1/4)
```

---

dtWV   *Noncentral t Distribution Density by W.V.*

Description

Compute the density function $f(x)$ of the t distribution with df degrees of freedom and non-centrality parameter ncp, according to Wolfgang Viechtbauer's proposal in 2002.

Usage

```
dtWV(x, df, ncp = 0, log = FALSE)
```

Arguments

x	numeric vector.
df	degrees of freedom (> 0, maybe non-integer). df = Inf is allowed.
ncp	non-centrality parameter δ; If omitted, use the central t distribution.
log	logical; if TRUE, $log(f(x))$ is returned instead of $f(x)$.

Details

The formula used is "asymptotic": Resnikoff and Lieberman (1957), p.1 and p.25ff, proposed to use recursive polynomials for (*integer !*) degrees of freedom $f = 1, 2, \ldots, 20$, and then, for df$=$ $f > 20$, use the asymptotic approximation which Wolfgang Viechtbauer proposed as a first version of a non-central t density for R (when dt() did not yet have an ncp argument).

Value

numeric vector of density values, properly recycled in (x,df,ncp).

Author(s)

Wolfgang Viechtbauer (2002) post to R-help (https://stat.ethz.ch/pipermail/r-help/2002-October/026044.html), and Martin Maechler (log argument; tweaks, notably recycling).

References

Resnikoff, George J. and Lieberman, Gerald J. (1957) *Tables of the non-central t-distribution*; Technical report no. 32 (LIE ONR 32), April 1, 1957; Applied Math. and Stat. Lab., Stanford University. https://statistics.stanford.edu/research/tables-non-central-t-distribution-density-function-cum

See Also

dt, R's (C level) implementation of the (non-central) t density.

Examples

```
tt <- seq(0, 10, len = 21)
ncp <- seq(0, 6, len = 31)
dt3R  <- outer(tt, ncp, dt  , df = 3)
dt3WV <- outer(tt, ncp, dtWV, df = 3)
all.equal(dt3R, dt3WV) # rel.err 0.00063
dt25R  <- outer(tt, ncp, dt  , df = 25)
dt25WV <- outer(tt, ncp, dtWV, df = 25)
all.equal(dt25R, dt25WV) # rel.err 1.1e-5

x <- -10:700
fx  <- dt  (x, df = 22, ncp =100)
lfx <- dt  (x, df = 22, ncp =100, log=TRUE)
lfV <- dtWV(x, df = 22, ncp =100, log=TRUE)

head(lfx, 20) # shows that R's dt(*, log=TRUE) implementation is "quite suboptimal"

## graphics
opa <- par(no.readonly=TRUE)
par(mar=.1+c(5,4,4,3), mgp = c(2, .8,0))
plot(fx ~ x, type="l")
par(new=TRUE) ; cc <- c("red", adjustcolor("orange", 0.4))
plot(lfx ~ x, type = "o", pch=".", col=cc[1], cex=2, ann=FALSE, yaxt="n")
sfsmisc::eaxis(4, col=cc[1], col.axis=cc[1], small.args = list(col=cc[1]))
lines(x, lfV, col=cc[2], lwd=3)
dtt1 <- "     dt"; dtt2 <- "(x, df=22, ncp=100"; dttL <- paste0(dtt2,", log=TRUE)")
legend("right", c(paste0(dtt1,dtt2,")")), paste0(c(dtt1,"dtWV"), dttL)),
       lty=1, lwd=c(1,1,3), col=c("black", cc), bty = "n")
par(opa) # reset
```

---

format01prec	*Format Numbers in [0,1] with "Precise" Result*

Description

Format numbers in [0,1] with "precise" result, notably using "1-.." if needed.

Usage

```
format01prec(x, digits = getOption("digits"), width = digits + 2,
             eps = 1e-06, ...,
             FUN = function(x, ...) formatC(x, flag = "-", ...))
```

Arguments

x	numbers in [0,1]; (still works if not)
digits	number of digits to use; is used as FUN(*,digits = digits) or FUN(*,digits = digits -5) depending on x or eps.
width	desired width (of strings in characters), is used as FUN(*,width = width) or FUN(*,width = width -2) depending on x or eps.
eps	small positive number: Use '1-' for those x which are in $(1 - eps, 1]$. The author has claimed in the last millennium that (the default) 1e-6 is *optimal*.

| ... | optional further arguments passed to FUN(x,digits,width,...). |
| FUN | a function used for format()ing; must accept both a digits and width argument. |

Value

a character vector of the same length as x.

Author(s)

Martin Maechler, 14 May 1997

See Also

formatC, format.pval.

Examples

```
## Show that format01prec()  does reveal more precision :
cbind(format      (1 - 2^-(16:24)),
      format01prec(1 - 2^-(16:24)))

## a bit more variety
e <- c(2^seq(-24,0, by=2), 10^-(7:1))
ee <- sort(unique(c(e, 1-e)))
noquote(ff <- format01prec(ee))
data.frame(ee, format01prec = ff)
```

---

lbeta   *(Log) Beta Approximations*

Description

Compute log(beta(a,b)) in a simple (fast) or asymptotic way.

Usage

```
lbetaM  (a, b, k.max = 5, give.all = FALSE)
lbeta_asy(a, b, k.max = 5, give.all = FALSE)
lbetaMM  (a, b, cutAsy = 1e-2, verbose = FALSE)

 betaI(a, n)
lbetaI(a, n)

logQab_asy(a, b, k.max = 5, give.all = FALSE)
Qab_terms(a, k)
```

Arguments

a, b, n	the Beta parameters, see [beta](); n must be a positive integer and "small".
k.max	..
give.all	[logical]() ..
cutAsy	cutoff value from where to switch to asymptotic formula.
verbose	logical (or integer) indicating if and how much monitoring information should be printed to the console.
k	the number of terms in the series expansion of Qab_terms(), currently must be in $\{0, 1, .., 5\}$.

Details

All lbeta*() functions compute log([beta](a,b)).

We use $Qab = Qab(a, b)$ for

$$Q_{a,b} := \frac{\Gamma(a+b)}{\Gamma(b)},$$

which is numerically challenging when b becomes large compared to a, or $a \ll b$.

With the beta function

$$B(a,b) = \frac{\Gamma(a)\Gamma(b)}{\Gamma(a+b)} = \frac{\Gamma(a)}{Qab},$$

and hence

$$\log B(a,b) = \log\Gamma(a) + \log\Gamma(b) - \log\Gamma(a+b) = \log\Gamma(a) - \log Qab,$$

or in R, lBeta(a,b) := lgamma(a) -logQab(a,b).

Indeed, typically everything has to be computed in log scale, as both $\Gamma(b)$ and $\Gamma(a+b)$ would overflow numerically for large b. Consequently, we use logQab*(), and for the large b case logQab_asy() specifically,

$$\texttt{logQab(a,b)} := \log(Qab(a,b)).$$

Note this is related to trying to get asymptotic formula for Γ ratios, notably formula (6.1.47) in Abramowitz and Stegun.

Note how this is related to computing [qbeta]()() in boundary cases, and see [algdiv]()() 'Details' about this.

We also have a vignette about this, but really the problem has been adressed pragmatically by the authors of TOMS 708, see the 'References' in [pbeta](), by their routine [algdiv]()() which also is available in our package **DPQ**.

Value

a fast or simple (approximate) computation of [lbeta](a,b).

Author(s)

Martin Maechler

References

Abramowitz, M. and Stegun, I. A. (1972) *Handbook of Mathematical Functions*. New York: Dover. https://en.wikipedia.org/wiki/Abramowitz_and_Stegun provides links to the full text which is in public domain.

Formula (6.1.47), p.257

See Also

R's beta function; algdiv().

Examples

```
## TODO
```

---

lgamma1p   *Accurate* log(gamma(a+1))

Description

Compute
$$l\Gamma_1(a) := \log\Gamma(a+1) = \log(a \cdot \Gamma(a)) = \log a + \log\Gamma(a),$$

which is "in principle" the same as log(gamma(a+1)) or lgamma(a+1), accurately also for (very) small a $(0 < a < 0.5)$.

Usage

```
lgamma1p (a, tol_logcf = 1e-14)
lgamma1p.(a, cutoff.a = 1e-6, k = 3)
lgamma1p_series(x, k)
```

Arguments

a, x	a numeric vector.
tol_logcf	for lgamma1p(): a non-negative number ...
cutoff.a	for lgamma1p.(): a positive number indicating the cutoff to switch from ...
k	an integer, the number of terms in the series expansion used internally.

Details

lgamma1p() is an R translation of the function (in Fortran) in Didonato and Morris (1992) which uses a 40-degree polynomial approximation.

lgamma1p_series(x,k) is Taylor series approximation of order k, (derived via Maple), which is $-\gamma x + \pi^2 x^2/12 + O(x^3)$, where γ is Euler's constant 0.5772156649....

Value

a numeric vector with the same attributes as a.

Author(s)

Morten Welinder (C code of Jan 2005, see R's bug issue PR#7307) for `lgamma1p()`.

Martin Maechler, notably for `lgamma1p_series()` which works with package **Rmpfr** but otherwise may be *much* less accurate than Morten's 40 term series!

References

Didonato, A. and Morris, A., Jr, (1992) Algorithm 708: Significant digit computation of the incomplete beta function ratios. *ACM Transactions on Mathematical Software*, **18**, 360–373; see also `pbeta`.

See Also

`log1pmx`, `log1p`, `pbeta`.

Examples

```
curve(-log(x*gamma(x)), 1e-30, .8, log="xy", col="gray50", lwd = 3,
      axes = FALSE, ylim = c(1e-30,1))
sfsmisc::eaxis(1); sfsmisc::eaxis(2)
at <- 10^(1-4*(0:8))
abline(h = at, v = at, col = "lightgray", lty = "dotted")
curve(-lgamma( 1+x), add=TRUE, col="red2", lwd=1/2)# underflows even earlier
curve(-lgamma1p (x), add=TRUE, col="blue")
curve(-lgamma1p.(x), add=TRUE, col=adjustcolor("forest green",1/4),
      lwd = 5, lty = 2)
for(k in 1:7)
  curve(-lgamma1p_series(x, k=k), add=TRUE, col=paste0("gray",30+k*8), lty = 3)
```

---

log1mexp   *Compute f(a) = log(1 - exp(-a)) Numerically Optimally*

Description

Compute f(a) = log(1 - exp(-a)) quickly and numerically accurately.

Usage

```
log1mexp(x)
```

Arguments

x   numeric vector of positive values.

Author(s)

Martin Maechler

References

Martin Mächler (2012). Accurately Computing $\log(1 - \exp(-|a|))$; https://CRAN.R-project.org/package=Rmpfr/vignettes/log1mexp-note.pdf.

See Also

The log1mexp() function in CRAN package **copula**, and the corresponding vignette (in the 'References').

---

log1pmx	*Accurate* log(1+x) - x

Description

Compute

$$\log(1 + x) - x$$

accurately also for small x, i.e., $|x| \ll 1$.

Usage

```
log1pmx(x, tol_logcf = 1e-14)
```

Arguments

x	numeric vector with values $x > -1$.
tol_logcf	a non-negative number indicating the tolerance (maximal relative error) for the auxiliary logcf() function.

Details

In order to provide full accuracy, the computations happens differently in three regions for x,

$$m_l = -0.79149064$$

is the first cutpoint,

$x < ml$ **or** $x > 1$**:** use log1pmx(x) := log1p(x) -x,

$|x| < 0.01$**:** use $t((((2/9 * y + 2/7)y + 2/5)y + 2/3)y - x)$,

$x \in [ml, 1]$**, and** $|x| >= 0.01$**:** use $t(2ylogcf(y, 3, 2) - x)$,

where $t := \frac{x}{2+x}$, and $y := t^2$.

Note that the formulas based on t are based on the (fast converging) formula

$$\log(1 + x) = 2 \left(r + \frac{r^3}{3} + \frac{r^5}{5} + \dots \right),$$

where $r := x/(x + 2)$, see the reference.

Value

a numeric vector (with the same attributes as x).

Author(s)

A translation of Morten Welinder's C code of Jan 2005, see R's bug issue PR#7307.

References

Abramowitz, M. and Stegun, I. A. (1972) *Handbook of Mathematical Functions*. New York: Dover. https://en.wikipedia.org/wiki/Abramowitz_and_Stegun provides links to the full text which is in public domain.
Formula (4.1.29), p.68.

See Also

logcf, the auxiliary function, lgamma1p which calls log1pmx, log1p

Examples

```
l1x <- curve(log1pmx, -.9999, 7, n=1001)
abline(h=0, v=-1:0, lty=3)
l1xz  <- curve(log1pmx, -.1, .1, n=1001); abline(h=0, v=0, lty=3)
l1xz2 <- curve(log1pmx, -.01, .01, n=1001); abline(h=0, v=0, lty=3)
l1xz3 <- curve(-log1pmx(x), -.002, .002, n=2001, log="y", yaxt="n")
sfsmisc::eaxis(2); abline(v=0, lty=3)
```

---

logcf   *Continued Fraction Approximation of Log-Related Series*

Description

Compute a continued fraction approximation to the series (infinite sum)

$$\sum_{k=0}^{\infty} \frac{x^k}{i + k \cdot d} = \frac{1}{i} + \frac{x}{i+d} + \frac{x^2}{i+2*d} + \frac{x^3}{i+3*d} + \ldots$$

Needed as auxiliary function in log1pmx() and lgamma1p().

Usage

```
logcf(x, i, d, eps, maxit = 10000)
```

Arguments

x	numeric vector
i	positive numeric
d	non-negative numeric
eps	positive number, the convergence tolerance.
maxit	a positive integer, the maximal number of iterations or terms in the truncated series used.

Value

a numeric vector with the same attributes as x.

Note

Rescaling is done by (namespace hidden) "global" scalefactor ....

Author(s)

Martin Maechler, based on ....

See Also

lgamma1p, log1pmx, and pbeta, whose prinicipal algorithm has evolved from TOMS 708.

Examples

```
l32 <- curve(logcf(x, 3,2, eps=1e-7), -3, 1)
abline(h=0,v=1, lty=3, col="gray50")
plot(y~x, l32, log="y", type = "o", main = "logcf(*, 3,2)  in log-scale")
```

---

logspace.add   *Logspace Arithmetix – Addition and Subtraction*

Description

Compute the log(arithm) of a sum (or difference) from the log of terms without causing overflows and without throwing away large handfuls of accuracy.

logspace.add(lx, ly):=

$$\log(\exp(lx) + \exp(ly))$$

logspace.sub(lx, ly):=

$$\log(\exp(lx) - \exp(ly))$$

Usage

```
logspace.add(lx, ly)
logspace.sub(lx, ly)
```

Arguments

lx, ly   numeric vectors, typically of the same length, but will be recycled to common
                length as with other R arithmetic.

Value

a numeric vector of the same length as x+y.

Note

This is really from R's C source code for pgamma(), i.e., '<R>/src/nmath/pgamma.c'

The function definitions are very simple, logspace.sub() using log1mexp().

Author(s)

Morten Welinder (for R's pgamma()); Martin Maechler

See Also

lsum, lssum; then pgamma()

Examples

```
set.seed(12)
ly <- rnorm(100, sd= 50)
lx <- ly + abs(rnorm(100, sd=100))  # lx - ly must be positive for *.sub()
stopifnot(exprs = {
   all.equal(logspace.add(lx,ly),
              log(exp(lx) + exp(ly)), tol=1e-14)
   all.equal(logspace.sub(lx,ly),
              log(exp(lx) - exp(ly)), tol=1e-14)
})
```

---

lssum   *Compute Logarithm of a Sum with Signed Large Summands*

Description

Properly compute $\log(x_1 + \ldots + x_n)$ for given log absolute values lxabs = $log(|x_1|), .., log(|x_n|)$ and corresponding signs signs = $sign(x_1), .., sign(x_n)$. Here, x_i is of arbitrary sign.

Notably this works in many cases where the direct sum would have summands that had overflown to +Inf or underflown to -Inf.

This is a (simpler, vector-only) version of copula:::lssum() (CRAN package **copula**).

Note that the *precision* is often not the problem for the direct summation, as R's sum() internally uses "long double" precision on most platforms.

Usage

```
lssum(lxabs, signs, l.off = max(lxabs), strict = TRUE)
```

Arguments

lxabs	n-vector of values $\log(|x_1|), \ldots, \log(|x_n|)$.
signs	corresponding signs $sign(x_1), \ldots, sign(x_n)$.
l.off	the offset to substract and re-add; ideally in the order of max(.).
strict	logical indicating if the function should stop on some negative sums.

Value

$$log(x_1 + .. + x_n) == log(sum(x)) = log(sum(sign(x) * |x|)) == log(sum(sign(x) * exp(log(|x|)))) == log(exp(log$$

Author(s)

Marius Hofert and Martin Maechler (for package **copula**).

See Also

lsum() which computes an exponential sum in log scale with*out* signs.

Examples

```
rSamp <- function(n, lmean, lsd = 1/4, roundN = 16) {
  lax <- sort((1+1e-14*rnorm(n))*round(roundN*rnorm(n, m = lmean, sd = lsd))/roundN)
  sx <- rep_len(c(-1,1), n)
  list(lax=lax, sx=sx, x = sx*exp(lax))
}

set.seed(101)
L1 <- rSamp(1000, lmean = 700) # here, lssum() is not needed (no under-/overflow)
summary(as.data.frame(L1))
ax <- exp(lax <- L1$lax)
hist(lax); rug(lax)
hist( ax); rug( ax)
sx <- L1$sx
table(sx)
(lsSimple <- log(sum(L1$x)))            # 700.0373
(lsS <- lssum(lxabs = lax, signs = sx))# ditto
lsS - lsSimple # even exactly zero (in 64b Fedora 30 Linux which has nice 'long double')
stopifnot(all.equal(700.037327351478, lsS, tol=1e-14), all.equal(lsS, lsSimple))

L2 <- within(L1, { lax <- lax + 10; x <- sx*exp(lax) }) ; summary(L2$x) # some -Inf, +Inf
(lsSimpl2 <- log(sum(L2$ x)))                    # NaN
(lsS2 <- lssum(lxabs = L2$ lax, signs = L2$ sx)) # 710.0373
stopifnot(all.equal(lsS2, lsS + 10, tol = 1e-14))
```

---

lsum   *Properly Compute the Logarithm of a Sum*

Description

Simple vector version of copula:::lsum() (CRAN package **copula**).

Properly compute $\log(x_1 + \ldots + x_n)$. for given $log(x_1), .., log(x_n)$. Here, $x_i > 0$ for all i.

Usage

```
lsum(lx, l.off = max(lx))
```

Arguments

lx	n-vector of values log(x_1),..,log(x_n).
l.off	the offset to substract and re-add; ideally in the order of the maximum of each column.

Value

$$log(x_1+..+x_n) = log(sum(x)) = log(sum(exp(log(x)))) == log(exp(log(x_max))*sum(exp(log(x)-log(x_max))$$

Author(s)

Originally, via paired programming: Marius Hofert and Martin Maechler.

See Also

[lssum](#)() which computes a sum in log scale with specified (typically alternating) signs.

Examples

```
## The "naive" version :
lsum0 <- function(lx) log(sum(exp(lx)))

lx1 <- 10*(-80:70) # is easy
lx2 <- 600:750     # lsum0() not ok [could work with rescaling]
lx3 <- -(750:900)  # lsum0() = -Inf - not good enough
m3 <- cbind(lx1,lx2,lx3)
lx6 <- lx5 <- lx4 <- lx3
lx4[149:151] <- -Inf ## = log(0)
lx5[150] <- Inf
lx6[1] <- NA_real_
m6 <- cbind(m3,lx4,lx5,lx6)
stopifnot(exprs = {
  all.equal(lsum(lx1), lsum0(lx1))
  all.equal((ls1 <- lsum(lx1)),  700.000045400960403, tol=8e-16)
  all.equal((ls2 <- lsum(lx2)),  750.458675145387133, tol=8e-16)
  all.equal((ls3 <- lsum(lx3)), -749.541324854612867, tol=8e-16)
  ## identical: matrix-version <==> vector versions
  identical(lsum(lx4), ls3)
  identical(lsum(lx4), lsum(head(lx4, -3))) # the last three were -Inf
  identical(lsum(lx5), Inf)
  identical(lsum(lx6), lx6[1])
  identical((lm3 <- apply(m3, 2, lsum)), c(lx1=ls1, lx2=ls2, lx3=ls3))
  identical(apply(m6, 2, lsum), c(lm3, lx4=ls3, lx5=Inf, lx6=lx6[1]))
})
```

---

newton | *Simple R level Newton Algorithm, Mostly for Didactical Reasons*

Description

Given the function G() and its derivative g(), newton() uses the Newton method, starting at x0, to find a point xp at which G is zero. G() and g() may each depend on the same parameter (vector) z.

Convergence typically happens when the stepsize becomes smaller than eps.

keepAll = TRUE to also get the vectors of consecutive values of x and G(x, z);

Usage

```
newton(x0, G, g, z,
       xMin = -Inf, xMax = Inf, warnRng = TRUE,
       dxMax = 1000, eps = 0.0001, maxiter = 1000L,
       warnIter = missing(maxiter) || maxiter >= 10L,
       keepAll = NA)
```

Arguments

x0	numeric start value.
G, g	must be [functions](#), mathematically of their first argument, but they can accept parameters; g() must be the derivative of G.
z	parameter vector for $G()$ and $g()$, to be kept fixed.
xMin, xMax	numbers defining the allowed range for x during the iterations; e.g., useful to set to 0 and 1 during quantile search.
warnRng	[logical](#) specifying if a [warning](#) should be signalled when start value x0 is outside [xMin,xMax] and hence will be changed to one of the boundary values.
dxMax	maximal step size in x-space. (The default 1000 is quite arbitrary, do set a good maximal step size yourself!)
eps	positive number, the *absolute* convergence tolerance.
maxiter	positive integer, specifying the maximal number of Newton iterations.
warnIter	logical specifying if a [warning](#) should be signalled when the algorithm has not converged in maxiter iterations.
keepAll	logical specifying if the full sequence of x- and G(x,*) values should be kept and returned:
	NA, the default: newton returns a small list of final "data", with 4 components x= $x*$, G= $G(x*, z)$, it, and converged.
	TRUE: returns an extended [list](#), in addition containing the vectors x.vec and G.vec.
	FALSE: returns only the $x*$ value.

Details

Because of the quadrativc convergence at the end of the Newton algorithm, often x^* satisfies approximately $|G(x*, z)| < eps^2$.

newton() can be used to compute the quantile function of a distribution, if you have a good starting value, and provide the cumulative probability and density functions as R functions G and g respectively.

Value

The result always contains the final x-value $x*$, and typically some information about convergence, depending on the value of keepAll, see above:

x	the optimal x^* value (a number).
G	the function value $G(x*, z)$, typically very close to zero.
it	the integer number of iterations used.
convergence	logical indicating if the Newton algorithm converged within maxiter iterations.
x.vec	the full vector of x values, $\{x0, \dots, x^*\}$.
G.vec	the vector of function values (typically tending to zero), i.e., G(x.vec,.) (even when G(x,.) would not vectorize).

Author(s)

Martin Maechler, ca. 2004

References

Newton's Method on Wikipedia, https://en.wikipedia.org/wiki/Newton%27s_method.

See Also

`uniroot()` is much more sophisticated, works without derivatives and is generally faster than `newton()`.

`newton(.)` is currently crucially used (only) in our function `qchisqN()`.

Examples

```
## The most simple non-trivial case :  Computing SQRT(a)
  G <- function(x, a) x^2 - a
  g <- function(x, a) 2*x

  newton(1, G, g, z = 4  ) # z = a -- converges immediately
  newton(1, G, g, z = 400) # bad start, needs longer to converge

## More interesting, and related to non-central (chisq, e.t.) computations:
## When is  x * log(x) < B,  i.e., the inverse function of G = x*log(x) :
xlx <- function(x, B) x*log(x) - B
dxlx <- function(x, B) log(x) + 1

Nxlx <- function(B) newton(B, G=xlx, g=dxlx, z=B, maxiter=Inf)$x
N1   <- function(B) newton(B, G=xlx, g=dxlx, z=B, maxiter = 1)$x
N2   <- function(B) newton(B, G=xlx, g=dxlx, z=B, maxiter = 2)$x

Bs <- c(outer(c(1,2,5), 10^(0:4)))
plot (Bs, vapply(Bs, Nxlx, pi), type = "l", log ="xy")
lines(Bs, vapply(Bs, N1  , pi), col = 2, lwd = 2, lty = 2)
lines(Bs, vapply(Bs, N2  , pi), col = 3, lwd = 3, lty = 3)

BL <- c(outer(c(1,2,5), 10^(0:6)))
plot (BL, vapply(BL, Nxlx, pi), type = "l", log ="xy")
lines(BL, BL, col="green2", lty=3)
lines(BL, vapply(BL, N1  , pi), col = 2, lwd = 2, lty = 2)
lines(BL, vapply(BL, N2  , pi), col = 3, lwd = 3, lty = 3)
## Better starting value from an approximate 1 step Newton:
iL1 <- function(B) 2*B / (log(B) + 1)
lines(BL, iL1(BL), lty=4, col="gray20") ## really better ==> use it as start

Nxlx <- function(B) newton(iL1(B), G=xlx, g=dxlx, z=B, maxiter=Inf)$x
N1   <- function(B) newton(iL1(B), G=xlx, g=dxlx, z=B, maxiter = 1)$x
N2   <- function(B) newton(iL1(B), G=xlx, g=dxlx, z=B, maxiter = 2)$x

plot (BL, vapply(BL, Nxlx, pi), type = "o", log ="xy")
lines(BL, iL1(BL),  lty=4, col="gray20")
lines(BL, vapply(BL, N1  , pi), type = "o", col = 2, lwd = 2, lty = 2)
lines(BL, vapply(BL, N2  , pi), type = "o", col = 3, lwd = 2, lty = 3)
## Manual 2-step Newton
iL2 <- function(B) { lB <- log(B) ; B*(lB+1) / (lB * (lB - log(lB) + 1)) }
lines(BL, iL2(BL), col = adjustcolor("sky blue", 0.6), lwd=6)
##==>  iL2() is very close to true curve
## relative error:
iLtrue <- vapply(BL, Nxlx, pi)
cbind(BL, iLtrue, iL2=iL2(BL), relErL2 = 1-iL2(BL)/iLtrue)
```

```
## absolute error (in log-log scale; always positive!):
plot(BL, iL2(BL) - iLtrue, type = "o", log="xy", axes=FALSE)
if(requireNamespace("sfsmisc")) {
  sfsmisc::eaxis(1)
  sfsmisc::eaxis(2, sub10=2)
} else {
  cat("no 'sfsmisc' package; maybe  install.packages(\"sfsmisc\")  ?\n")
  axis(1); axis(2)
}
## 1 step from iL2()  seems quite good:
B. <- BL[-1] # starts at 2
NL2 <- lapply(B., function(B) newton(iL2(B), G=xlx, g=dxlx, z=B, maxiter=1))
str(NL2)
iL3 <- sapply(NL2, `[[`, "x")
cbind(B., iLtrue[-1], iL2=iL2(B.), iL3, relE.3 = 1- iL3/iLtrue[-1])
x. <- iL2(B.)
all.equal(iL3, x. - xlx(x., B.) / dxlx(x.)) ## 7.471802e-8
## Algebraic simplification of one newton step :
all.equal((x.+B.)/(log(x.)+1), x. - xlx(x., B.) / dxlx(x.), tol = 4e-16)
iN1 <- function(x, B) (x+B) / (log(x) + 1)
B <- 12345
iN1(iN1(iN1(B, B),B),B)
Nxlx(B)
```

---

numer-utils   *Numerical Utilities - Functions, Constants*

Description

The **DPQ** package provides some numeric constants used in some of its distribution computations.

all_mpfr() and any_mpfr() return [TRUE](#) iff all (or 'any', respectively) of their arguments inherit from [class](#) "mpfr" (from package **Rmpfr**).

logr(x,a) computes [log](#)(x / (x + a)) in a numerically stable way.

Usage

```
## Numeric Constants : % mostly in   ../R/beta-fns.R
M_LN2       # = log(2) = 0.693....
M_cutoff    # := If |x| > |k| * M_cutoff, then  log[ exp(-x) * k^x ]  =~=  -x
            #  = 3196577161300663808 ~= 3.2e+18
M_minExp    # = log(2) * .Machine$double.min.exp # ~= -708.396..
G_half      # = sqrt(pi) = Gamma( 1/2 )

## Functions :
all_mpfr(...)
any_mpfr(...)
logr(x, a)   # == log(x / (x + a)) -- but numerically smart; x > 0, a >= 0 > -x
```

Arguments

...	numeric or "mpfr" numeric vectors.
x, a	number-like ..

Details

all_mpfr(),

all_mpfr(): test if [all](#) or [any](#) of their arguments or of class "mpfr" (from package **Rmpfr**). The arguments are evaluated only until the result is determined, see the example.

logr() computes $\log(x/(x + a))$ in a numerically stable way.

Value

The numeric constant in the first case; a numeric (or "mpfr") vector of appropriate size in the 2nd case.

Author(s)

Martin Maechler

See Also

[.Machine](#)

Examples

```
(Ms <- ls("package:DPQ", pattern = "^M"))
lapply(Ms, function(nm) { cat(nm,": "); print(get(nm)) }) -> .tmp
```

---

pbetaRv1   *Pure R Implementation of Old pbeta()*

Description

pbetaRv1() is an implementation of the original ("version 1" [pbeta](#)() function in R (versions <= 2.2.x), before we started using TOMS 708 bratio() instead, see that help page also for references.

pbetaRv1() is basically a manual translation from C to R of the underlying pbeta_raw() C function, see in R's source tree at https://svn.r-project.org/R/branches/R-2-2-patches/src/nmath/pbeta.c

For consistency within R, we are using R's argument names (q,shape1,shape2) instead of C code's (x,pin,qin).

It is only for the *central* beta distribution.

Usage

```
pbetaRv1(q, shape1, shape2, lower.tail = TRUE,
         eps = 0.5 * .Machine$double.eps,
         sml = .Machine$double.xmin,
         verbose = 0)
```

Arguments

q, shape1, shape2

           non-negative numbers, q in $[0, 1]$, see [pbeta](#).

lower.tail   indicating if $F(q; *)$ should be returned or the upper tail probability $1 - F(q)$.

eps   the tolerance used to determine congerence. eps has been hard coded in C code to `0.5 * .Machine$double.eps` which is equal to 2^{-53} or `1.110223e-16`.

sml   the smallest positive number on the typical platform. The default `.Machine$double.xmin` is hard coded in the C code (as `DBL_MIN`), and this is equal to 2^{-1022} or `2.225074e-308` on all current platforms.

verbose   integer indicating the amount of verbosity of diagnostic output, 0 means no output, 1 more, etc.

Value

a number.

Note

The C code contains
This routine is a translation into C of a Fortran subroutine by W. Fullerton of Los Alamos Scientific Laboratory.

Author(s)

Martin Maechler

References

(From the C code:)

Nancy E. Bosten and E.L. Battiste (1974). Remark on Algorithm 179 (S14): Incomplete Beta Ratio. *Communications of the ACM*, **17**(3), 156–7.

See Also

[pbeta](#).

Examples

```
all.equal(pbetaRv1(1/4, 2, 3),
          pbeta   (1/4, 2, 3))
set.seed(101)

N <- 1000
x <- sample.int(7, N, replace=TRUE) / 8
a <-   rlnorm(N)
b <- 5*rlnorm(N)
pbt <- pbeta(x, a, b)
for(i in 1:N) {
   stopifnot(all.equal(pbetaRv1(x[i], a[i], b[i]), pbt[i]))
   cat(".", if(i %% 20 == 0) paste0(i, "\n"))
}
```

---

pl2curves   *Plot 2 Noncentral Distribution Curves for Visual Comparison*

Description

Plot two noncentral (chi-squared or *t* or ..) distribution curves for visual comparison.

Usage

```
pl2curves(fun1, fun2, df, ncp, log = FALSE,
          from = 0, to = 2 * ncp, p.log = "", n = 2001,
          leg = TRUE, col2 = 2, lwd2 = 2, lty2 = 3, ...)
```

Arguments

fun1, fun2	[function]( )s, both to be used via [curve]( ), and called with the same 4 arguments, (.,df,ncp,log) (the name of the first argument is not specified).
df, ncp, log	parameters to be passed and used in both functions, which hence typically are non-central chi-squared or t density, probability or quantile functions.
from, to	numbers determining the x-range, passed to [curve]( ).
p.log	string, passed as [curve](....,log = log.p).
n	the number of evaluation points, passed to [curve]( ).
leg	logical specifying if a [legend]( ) should be drawn.
col2, lwd2, lty2	
	color, line width and line type for the second curve. (The first curve uses defaults for these graphical properties.)
...	further arguments passed to *first* [curve](..) call.

Value

TODO: inivisible return both curve() results, i.e., (x,y1, y2), possibly as data frame

Author(s)

Martin Maechler

See Also

[curve]( ), ..

Examples

```
p.dnchiBessel <- function(df, ncp, log=FALSE, from=0, to = 2*ncp, p.log="", ...)
{
    pl2curves(dnchisqBessel, dchisq, df=df, ncp=ncp, log=log,
              from=from, to=to, p.log=p.log, ...)
}

  ## TODO the p.dnchiB() examples >>>>>> ../tests/chisq-nonc-ex.R <<<
```

pnbeta   *Noncentral Beta Probabilities*

Description

pnbetaAppr2() and its inital version pnbetaAppr2v1() provide the "approximation 2" of Chattamvelli and Shanmugam(1997) to the noncentral Beta probability distribution.

pnbetaAS310() is an R level interface to a C translation (and "Rification") of the AS 310 Fortran implementation.

Usage

```
pnbetaAppr2(x, a, b, ncp = 0, lower.tail = TRUE, log.p = FALSE)
```

Arguments

x	numeric vector (of quantiles), typically from inside $[0, 1]$.
a, b	the shape parameters of Beta, aka as shape1 and shape2.
ncp	non-centrality parameter.
log.p	logical; if TRUE, probabilities p are given as log(p).
lower.tail	logical; if TRUE (default), probabilities are $P[X \leq x]$, otherwise, $P[X > x]$.

Value

a numeric vector of (log) probabilities of the same length as x.

Note

The authors in the reference compare AS 310 with Frick(1990) and Lenth(1987) and state to be better than both. R's current (2019) noncentral beta implementation builds on Lenth(1987), with some amendments though; still, pnbetaAS310() may potentially be better, at least in certain corners of the 4-dimensional input space.

Author(s)

Martin Maechler; pnbetaAppr2() in Oct 2007.

References

Chattamvelli, R., and Shanmugam, R. (1997) Algorithm AS 310: Computing the Non-Central Beta Distribution Function. *Journal of the Royal Statistical Society. Series C (Applied Statistics)* **46**(1), 146–156, for "approximation 2" notably p.154;

doi: 10.1111/14679876.00055.

See Also

R's own pbeta.

Examples

```
## Same arguments as for Table 1 (p.151) of the reference
a <- 5*rep(1:3, each=3)
aargs <- cbind(a = a, b = a,
               ncp = rep(c(54, 140, 170), 3),
               x = 1e-4*c(8640, 9000, 9560, 8686, 9000, 9000, 8787, 9000, 9220))
aargs
pnbA2 <- apply(aargs, 1, function(aa) do.call(pnbetaAppr2, as.list(aa)))
pnA310<- apply(aargs, 1, function(aa) do.call(pnbetaAS310, as.list(aa)))
aar2 <- aargs; dimnames(aar2)[[2]] <- c(paste0("shape", 1:2), "ncp", "q")
pnbR  <- apply(aar2,  1, function(aa) do.call(pbeta, as.list(aa)))
range(relD2    <- 1 - pnbA2 /pnbR)
range(relD310 <- 1 - pnA310/pnbR)
cbind(aargs, pnbA2, pnA310, pnbR,
      relD2 = signif(relD2, 3), relD310 = signif(relD310, 3)) # <------> Table 1
stopifnot(abs(relD2)   < 0.009) # max is 0.006286
stopifnot(abs(relD310) < 1e-5 ) # max is 6.3732e-6

## Arguments as for Table 2 (p.152) of the reference :
aarg2 <- cbind(a = c( 10, 10, 15, 20, 20, 20, 30, 30),
               b = c( 20, 10,  5, 10, 30, 50, 20, 40),
               ncp=c(150,120, 80,110, 65,130, 80,130),
               x = c(868,900,880,850,660,720,720,800)/1000)
pnbA2 <- apply(aarg2, 1, function(aa) do.call(pnbetaAppr2, as.list(aa)))
pnA310<- apply(aarg2, 1, function(aa) do.call(pnbetaAS310, as.list(aa)))
aar2 <- aarg2; dimnames(aar2)[[2]] <- c(paste0("shape", 1:2), "ncp", "q")
pnbR  <- apply(aar2,  1, function(aa) do.call(pbeta, as.list(aa)))
range(relD2    <- 1 - pnbA2 /pnbR)
range(relD310 <- 1 - pnA310/pnbR)
cbind(aarg2, pnbA2, pnA310, pnbR,
      relD2 = signif(relD2, 3), relD310 = signif(relD310, 3)) # <------> Table 2
stopifnot(abs(relD2  ) < 0.006) # max is 0.00412
stopifnot(abs(relD310) < 1e-5 ) # max is 5.5953e-6

## Arguments as for Table 3 (p.152) of the reference :
aarg3 <- cbind(a = c( 10, 10, 10, 15, 10, 12, 30, 35),
               b = c(  5, 10, 30, 20,  5, 17, 30, 30),
               ncp=c( 20, 54, 80,120, 55, 64,140, 20),
               x = c(644,700,780,760,795,560,800,670)/1000)
pnbA3 <- apply(aarg3, 1, function(aa) do.call(pnbetaAppr2, as.list(aa)))
pnA310<- apply(aarg3, 1, function(aa) do.call(pnbetaAS310, as.list(aa)))
aar3 <- aarg3; dimnames(aar3)[[2]] <- c(paste0("shape", 1:2), "ncp", "q")
pnbR  <- apply(aar3,  1, function(aa) do.call(pbeta, as.list(aa)))
range(relD2    <- 1 - pnbA3 /pnbR)
range(relD310 <- 1 - pnA310/pnbR)
cbind(aarg3, pnbA3, pnA310, pnbR,
      relD2 = signif(relD2, 3), relD310 = signif(relD310, 3)) # <------> Table 3
stopifnot(abs(relD2  ) < 0.09) # max is 0.06337
stopifnot(abs(relD310) < 1e-4) # max is 3.898e-5
```

---

pnchi1sq   *(Probabilities of Non-Central Chisquared Distribution for Special*
                                Cases

Description

Computes probabilities for the non-central chi squared distribution, in special cases, currently for df = 1 and df = 3, using 'exact' formulas only involving the standard normal (Gaussian) cdf $\Phi()$ and its derivative $\phi()$, i.e., R's pnorm() and dnorm().

Usage

```
pnchi1sq(q, ncp = 0, lower.tail = TRUE, log.p = FALSE, epsS = .01)
pnchi3sq(q, ncp = 0, lower.tail = TRUE, log.p = FALSE, epsS = .04)
```

Arguments

q	number ('quantile', i.e., abscissa value.)
ncp	non-centrality parameter δ; ....
lower.tail, log.p	
	logical, see, e.g., pchisq().
epsS	small number, determining where to switch from the "small case" to the regular case, namely by defining small <-sqrt(q/ncp) <= epsS.

Details

In the "small case" (epsS above), the direct formulas suffer from cancellation, and we use Taylor series expansions in $s := \sqrt{q}$, which in turn use "probabilists"' Hermite polynomials $He_n(x)$.

The default values epsS have currently been determined by experiments as those in the 'Examples' below.

Value

a numeric vector "like" q+ncp, i.e., recycled to common length.

Author(s)

Martin Maechler, notably the Taylor approximations in the "small" cases.

References

Johnson et al.(1995), see 'References' in pnchisqPearson.

https://en.wikipedia.org/wiki/Hermite_polynomials

See Also

pchisq, the (simple and R-like) approximations, such as pnchisqPearson and the wienergerm approximations, pchisqW() etc.

Examples

```
qq <- seq(9500, 10500, length=1000)
m1 <- cbind(pch = pchisq  (qq, df=1, ncp = 10000),
            p1  = pnchi1sq(qq,       ncp = 10000))
matplot(qq, m1, type = "l"); abline(h=0:1, v=10000+1, lty=3)
all.equal(m1[,"p1"], m1[,"pch"], tol=0) # for now,  2.37e-12

m3 <- cbind(pch = pchisq  (qq, df=3, ncp = 10000),
```

```
                 p3 = pnchi3sq(qq,        ncp = 10000))
matplot(qq, m3, type = "l"); abline(h=0:1, v=10000+3, lty=3)
all.equal(m3[,"p3"], m3[,"pch"], tol=0) # for now,  1.88e-12

stopifnot(exprs = {
  all.equal(m1[,"p1"], m1[,"pch"], tol=1e-10)
  all.equal(m3[,"p3"], m3[,"pch"], tol=1e-10)
})

### Very small 'x' i.e., 'q' would lead to cancellation: -----------

##  df = 1 --------------

qS <- c(0, 2^seq(-40,4, by=1/16))
m1s <- cbind(pch = pchisq  (qS, df=1, ncp = 1)
           , p1.0= pnchi1sq(qS,        ncp = 1, epsS = 0)
           , p1.4= pnchi1sq(qS,        ncp = 1, epsS = 1e-4)
           , p1.3= pnchi1sq(qS,        ncp = 1, epsS = 1e-3)
           , p1.2= pnchi1sq(qS,        ncp = 1, epsS = 1e-2)
        )
cols <- adjustcolor(1:5, 1/2); lws <- seq(4,2, by = -1/2)
abl.leg <- function(x.leg = "topright", epsS = 10^-(4:2), legend = NULL)
{
    abline(h = .Machine$double.eps, v = epsS^2,
           lty = c(2,3,3,3), col= adjustcolor(1, 1/2))
    if(is.null(legend))
      legend <- c(quote(epsS == 0), as.expression(lapply(epsS,
                            function(K) substitute(epsS == KK,
                                                   list(KK = formatC(K, w=1))))))
    legend(x.leg, legend, lty=1:4, col=cols, lwd=lws, bty="n")
}
matplot(qS, m1s, type = "l", log="y" , col=cols, lwd=lws)
matplot(qS, m1s, type = "l", log="xy", col=cols, lwd=lws) ; abl.leg("right")
## ====  "Errors" =================================================
## Absolute: ------------------------
matplot(qS,     m1s[,1] - m1s[,-1] , type = "l", log="x" , col=cols, lwd=lws)
matplot(qS, abs(m1s[,1] - m1s[,-1]), type = "l", log="xy", col=cols, lwd=lws)
abl.leg("bottomright")
## Relative: ------------------------
matplot(qS,     1 - m1s[,-1]/m1s[,1] , type = "l", log="x",  col=cols, lwd=lws)
abl.leg()
matplot(qS, abs(1 - m1s[,-1]/m1s[,1]), type = "l", log="xy", col=cols, lwd=lws)
abl.leg()


##  df = 3 -------------- %% FIXME:  the 'small' case is clearly wrong <<<

qS <- c(0, 2^seq(-40,4, by=1/16))
ee <- c(1e-3, 1e-2, .04)
m3s <- cbind(pch = pchisq  (qS, df=3, ncp = 1)
           , p1.0= pnchi3sq(qS,        ncp = 1, epsS = 0)
           , p1.3= pnchi3sq(qS,        ncp = 1, epsS = ee[1])
           , p1.2= pnchi3sq(qS,        ncp = 1, epsS = ee[2])
           , p1.1= pnchi3sq(qS,        ncp = 1, epsS = ee[3])
        )
matplot(qS, m3s, type = "l", log="y" , col=cols, lwd=lws)
matplot(qS, m3s, type = "l", log="xy", col=cols, lwd=lws); abl.leg("right", ee)
```

```
## ====  "Errors" ===================================================
## Absolute: -----------------------
matplot(qS,     m3s[,1] - m3s[,-1] , type = "l", log="x" , col=cols, lwd=lws)
matplot(qS, abs(m3s[,1] - m3s[,-1]), type = "l", log="xy", col=cols, lwd=lws)
abl.leg("right", ee)
## Relative: -----------------------
matplot(qS,     1 - m3s[,-1]/m3s[,1] , type = "l", log="x",  col=cols, lwd=lws)
abl.leg(, ee)
matplot(qS, abs(1 - m3s[,-1]/m3s[,1]), type = "l", log="xy", col=cols, lwd=lws)
abl.leg(, ee)
```

---

pnchisqAppr   *(Approximate) Probabilities of Non-Central Chisquared Distribution*

Description

Compute (approximate) probabilities for the non-central chi squared distribution.

The non-central chi-squared distribution with df$= n$ degrees of freedom and non-centrality parameter ncp $= \lambda$ has density

$$f(x) = f_{n,\lambda}(x) = e^{-\lambda/2} \sum_{r=0}^{\infty} \frac{(\lambda/2)^r}{r!} f_{n+2r}(x)$$

for $x \geq 0$; for more, see R's help page for `pchisq`.

- R's own historical and current versions, but with more tuning parameters;

Historical relatively simple approximations listed in Johnson, Kotz, and Balakrishnan (1995):

- Patnaik(1949)'s approximation to the non-central via central chi-squared. Is also the formula 26.4.27 in Abramowitz & Stegun, p.942. Johnson et al mention that the approxmation error is $O(1/\sqrt(\lambda))$ *for* $\lambda \to \infty$.

- Pearson(1959) is using 3 moments instead of 2 as Patnaik (to approximate via a central chi-squared), and therefore better than Patnaik for the right tail; further (in Johnson et al.), the approxmation error is $O(1/\lambda)$ *for* $\lambda \to \infty$.

- Abdel-Aty(1954)'s "first approximation" based on Wilson-Hilferty via Gaussian (`pnorm`) probabilities, is partly *wrongly* cited in Johnson et al., p.463, eq.$(29.61a)$.

- Bol'shev and Kuznetzov (1963) concentrate on the case of **small** ncp λ and provide an "approximation" via *central* chi-squared with the same degrees of freedom df, but a modified q ('x'); the approximation has error $O(\lambda^3)$ for $\lambda \to 0$ and is from Johnson et al., p.465, eq.(29.62) and (29.63).

- Sankaran(1959, 1963) proposes several further approximations base on Gaussian probabilities, according to Johnson et al., p.463. pnchisqSankaran_d() implements its formula $(29.61d)$.

pnchisq(): an R implementation of R's own C pnchisq_raw(), but almost only up to Feb.27, 2004, long before the log.p=TRUE addition there, including *logspace arithmetic* in April 2014, its finish on 2015-09-01. Currently for historical reference only.

pnchisqV(): a `Vectorize`()d `pnchisq`.

pnchisqRC(): R's C implementation as of Aug.2019; but with many more options.

pnchisqIT: ....

pnchisqTerms: ....

pnchisqT93: pure R implementations of approximations when both q and ncp are large, by Temme(1993), from Johnson et al., p.467, formulas $(29.71a)$, and $(29.71b)$, using auxiliary functions pnchisqT93a() and pnchisqT93b() respectively, with adapted formulas for the log.p=TRUE cases.

pnchisq_ss(): ....

ss: ....

ss2: ....

ss2.: ....

Usage

```
pnchisq         (q, df, ncp = 0, lower.tail = TRUE,
                cutOffncp = 80, itSimple = 110, errmax = 1e-12, reltol = 1e-11,
                 maxit = 10* 10000, verbose = 0, xLrg.sigma = 5)
pnchisqV(x, ..., verbose = 0)

pnchisqRC        (q, df, ncp = 0, lower.tail = TRUE, log.p = FALSE,
                 no2nd.call = FALSE,
                 cutOffncp = 80, small.ncp.logspace = small.ncp.logspaceR2015,
                 itSimple = 110, errmax = 1e-12,
                reltol = 8 * .Machine$double.eps, epsS = reltol/2, maxit = 1e6,
                 verbose = FALSE)
pnchisqAbdelAty  (q, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
pnchisqBolKuz    (q, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
pnchisqPatnaik   (q, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
pnchisqPearson   (q, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
pnchisqSankaran_d(q, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
pnchisq_ss       (x, df, ncp = 0, lower.tail = TRUE, log.p = FALSE, i.max = 10000)
pnchisqTerms     (x, df, ncp,     lower.tail = TRUE, i.max = 1000)

pnchisqT93  (q, df, ncp, lower.tail = TRUE, log.p = FALSE, use.a = q > ncp)
pnchisqT93.a(q, df, ncp, lower.tail = TRUE, log.p = FALSE)
pnchisqT93.b(q, df, ncp, lower.tail = TRUE, log.p = FALSE)

ss  (x, df, ncp, i.max = 10000, useLv = !(expMin < -lambda && 1/lambda < expMax))
ss2 (x, df, ncp, i.max = 10000, eps = .Machine$double.eps)
ss2. (q, df, ncp = 0, errmax = 1e-12, reltol = 2 * .Machine$double.eps,
      maxit = 1e+05, eps = reltol, verbose = FALSE)
```

Arguments

x   numeric vector (of 'quantiles', i.e., abscissa values).

q   number ('quantile', i.e., abscissa value.)

df   degrees of freedom > 0, maybe non-integer.

ncp   non-centrality parameter δ; ....

lower.tail, log.p

             logical, see, e.g., [pchisq]().

`i.max`	number of terms in evaluation ...
`use.a`	[logical](#) vector for Temme pnchisqT93*() formulas, indicating to use formula 'a' over 'b'. The default is as recommended in the references, but they did not take into account `log.p = TRUE` situations.
`cutOffncp`	a positive number, the cutoff value for `ncp`...
`itSimple`	...
`errmax`	absolute error tolerance.
`reltol`	convergence tolerance for *rel*ative error.
`maxit`	maximal number of iterations.
`xLrg.sigma`	positive number ...
`no2nd.call`	logical indicating if a 2nd call is made to the internal function ....
`small.ncp.logspace`	
	logical vector or [function](#), indicating if the logspace computations for "small" ncp (defined to fulfill `ncp < cutOffncp` !).
`epsS`	small positive number, the convergence tolerance of the 'simple' iterations...
`verbose`	logical or integer specifying if or how much the algorithm progress should be monitored.
`...`	further arguments passed from pnchisqV() to pnchisq().
`useLv`	[logical](#) indicating if logarithmic scale should be used for λ computations.
`eps`	convergence tolerance, a positive number.

Details

pnchisq_ss() uses `si <-ss(x,df,..)` to get the series terms, and returns `2*`[dchisq](#)`(x,df = df +2) * sum(si$s)`.

ss() computes the terms needed for the expansion used in pnchisq_ss().

ss2() computes some simple "statistics" about ss(..).

Value

ss() returns a list with 3 components

`s`	the series
`i1`	location (in `s[]`) of the first change from 0 to positive.
`max`	(first) location of the maximal value in the series (i.e., [which.max](#)(s)).

Author(s)

Martin Maechler, from May 1999; starting from a post to the S-news mailing list by Ranjan Maitra (@ math.umbc.edu) who showed a version of our pchisqAppr.0() thanking Jim Stapleton for providing it.

References

Johnson, N.L., Kotz, S. and Balakrishnan, N. (1995) Continuous Univariate Distributions Vol~2, 2nd ed.; Wiley.
Chapter 29 *Noncentral χ^2-Distributions*; notably Section *8 Approximations*, p.461 ff.

Abramowitz, M. and Stegun, I. A. (1972) *Handbook of Mathematical Functions*. New York: Dover.
https://en.wikipedia.org/wiki/Abramowitz_and_Stegun

See Also

pchisq and the wienergerm approximations for it: pchisqW() etc.

r_pois() and its plot function, for an aspect of the series approximations we use in pnchisq_ss().

Examples

```
qq <- c(.001, .005, .01, .05, (1:9)/10, 2^seq(0, 10, by= 0.5))
pkg <- "package:DPQ"
pnchNms <- c(paste0("pchisq", c("V", "W", "W.", "W.R")),
             ls(pkg, pattern = "^pnchisq"))
pnchNms <- pnchNms[!grepl("Terms$", pnchNms)]
pnchF <- sapply(pnchNms, get, envir = as.environment(pkg))
str(pnchF)
for(ncp in c(0, 1/8, 1/2)) {
  cat("\n~~~~~~~~~~~~~\nncp: ", ncp,"\n======\n")
  pnF <- if(ncp == 0) pnchF[!grepl("chisqT93", pnchNms)] else pnchF
 print(sapply(pnF, function(F) Vectorize(F, names(formals(F))[[1]])(qq, df = 3, ncp=ncp)))
}

## A case where the non-central P[] should be improved :
## First, the central P[] which is close to exact -- choosing df=2 allows
## truly exact values: chi^2 = Exp(1) !
opal <- palette()
palette(c("black", "red", "green3", "blue", "cyan", "magenta", "gold3", "gray44"))
cR  <- curve(pchisq   (x, df=2,        lower.tail=FALSE, log.p=TRUE), 0, 4000, n=2001)
cRC <- curve(pnchisqRC(x, df=2, ncp=0, lower.tail=FALSE, log.p=TRUE),
             add=TRUE, col=adjustcolor(2,1/2), lwd=3, lty=2, n=2001)
cR0 <- curve(pchisq   (x, df=2, ncp=0, lower.tail=FALSE, log.p=TRUE),
             add=TRUE, col=adjustcolor(3,1/2), lwd=4,        n=2001)
## smart "named list" constructur :
list_ <- function(...)
   `names<-`(list(...), vapply(sys.call()[-1L], as.character, ""))
JKBfn <-list_(pnchisqPatnaik,
              pnchisqPearson,
              pnchisqAbdelAty,
              pnchisqBolKuz,
              pnchisqSankaran_d)
cl. <- setNames(adjustcolor(3+seq_along(JKBfn), 1/2), names(JKBfn))
lw. <- setNames(2+seq_along(JKBfn),                   names(JKBfn))
cR.JKB <- sapply(names(JKBfn), function(nmf) {
  curve(JKBfn[[nmf]](x, df=2, ncp=0, lower.tail=FALSE, log.p=TRUE),
        add=TRUE, col=cl.[[nmf]], lwd=lw.[[nmf]], lty=lw.[[nmf]], n=2001)
})
legend("bottomleft", c("pchisq", "pchisq.ncp=0", "pnchisqRC", names(JKBfn)),
       col=c(palette()[1], adjustcolor(2:3,1/2), cl.),
       lwd=c(1,3,4, lw.), lty=c(1,2,1, lw.))
palette(opal)# revert

all.equal(cRC, cR0, tol = 1e-15) # TRUE [for now]
## the problematic "jump" :
as.data.frame(cRC)[744:750,]
## verbose=TRUE  may reveal which branches of the algorithm are taken:
pnchisqRC(1500, df=2, ncp=0, lower.tail=FALSE, log.p=TRUE, verbose=TRUE) #
## |--> -Inf currently

## The *two*  principal cases (both lower.tail = {TRUE,FALSE} !), where
```

```
## "2nd call" happens *and* is currently beneficial :
dfs <- c(1:2, 5, 10, 20)
pL. <- pnchisqRC(.00001, df=dfs, ncp=0, log.p=TRUE, lower.tail=FALSE, verbose = TRUE)
pR. <- pnchisqRC(   100, df=dfs, ncp=0, log.p=TRUE,                    verbose = TRUE)
## R's own non-central version (specifying 'ncp'):
pL0 <- pchisq   (.00001, df=dfs, ncp=0, log.p=TRUE, lower.tail=FALSE)
pR0 <- pchisq   (   100, df=dfs, ncp=0, log.p=TRUE)
## R's *central* version, i.e., *not* specifying 'ncp' :
pL  <- pchisq   (.00001, df=dfs,        log.p=TRUE, lower.tail=FALSE)
pR  <- pchisq   (   100, df=dfs,        log.p=TRUE)
cbind(pL., pL, relEc = signif(1-pL./pL, 3), relE0 = signif(1-pL./pL0, 3))
cbind(pR., pR, relEc = signif(1-pR./pR, 3), relE0 = signif(1-pR./pR0, 3))
```

---

pnchisqWienergerm   *Wienergerm Approximations to (Non-Central) Chisquare Probabilities*

Description

Functions implementing the two Wiener germ approximations to [pchisq](), the (non-central) chisquare distribution, and to [qchisq]() its inverse, the quantile function.

These have been proposed by Penev and Raykov (2000) who also listed a Fortran implementation.

In order to use them in numeric boundary cases, Martin Maechler has improved the original formulas.

Auxiliary functions:

sW(): The $s()$ as in the Wienergerm approximation, but using Taylor expansion when needed, i.e., (x*ncp / df^2) << 1.

qs(): ...

z0(): ...

z.f(): ...

z.s(): ...

....

Usage

```
pchisqW. (q, df, ncp = 0, lower.tail = TRUE, log.p = FALSE,
          Fortran = TRUE, variant = c("s", "f"))
pchisqV  (q, df, ncp = 0, lower.tail = TRUE, log.p = FALSE,
          Fortran = TRUE, variant = c("s", "f"))
pchisqW  (q, df, ncp = 0, lower.tail = TRUE, log.p = FALSE, variant = c("s", "f"))
pchisqW.R(x, df, ncp = 0, lower.tail = TRUE, log.p = FALSE, variant = c("s", "f"),
          verbose = getOption("verbose"))

sW(x, df, ncp)
qs(x, df, ncp, f.s = sW(x, df, ncp), eps1 = 1/2, sMax = 1e+100)
z0(x, df, ncp)
z.f(x, df, ncp)
z.s(x, df, ncp, verbose = getOption("verbose"))
```

Arguments

q,x	vector of quantiles (main argument, see [pchisq](#)).
df	degrees of freedom (non-negative, but can be non-integer).
ncp	non-centrality parameter (non-negative).
lower.tail,log.p	
	[logical](#), see [pchisq](#).
variant	a [character](#) string, currently either "f" for the **f**irst or "s" for the **s**econd Wienergerm approximation in Penev \& Raykov (2000).
Fortran	logical specifying if the Fortran or the C version should be used.
verbose	logical (or integer) indicating if or how much diagnostic output should be printed to the console during the computations.
f.s	a number must be a "version" of $s(x, df, ncp)$.
eps1	for qs(): use direct approximation instead of h(1 -1/s) for s < eps1.
sMax	for qs(): cutoff to switch the $h(.)$ formula for s > sMax.

Details

....TODO... or write vignette

Value

all these functions return [numeric](#) vectors according to their arguments.

Note

The exact auxiliary function names etc, are still considered *provisional*; currently they are exported for easier documentation and use, but may well all disappear from the exported functions or even completely.

Author(s)

Martin Maechler, mostly end of Jan 2004

References

Penev, Spiridon and Raykov, Tenko (2000) A Wiener Germ approximation of the noncentral chi square distribution and of its quantiles. Computational Statistics **15**, 219–228. doi: [10.1007/s001800000029](#)

Dinges, H. (1989) Special cases of second order Wiener germ approximations. *Probability Theory and Related Fields*, **83**, 5–57.

See Also

[pchisq](#), and other approximations for it: [pnchisq](#)() etc.

Examples

```
## see  example(pnchisqAppr)   which looks at all of the pchisq() approximating functions
```

| pnt | *Non-central t Probability Distribution - Algorithms and Approximations* |

Description

Compute different approximations for the non-central t-Distribution cumulative probability distribution function.

Usage

```
pntR      (t, df, ncp, lower.tail = TRUE, log.p = FALSE,
                           itrmax = 1000, errmax = 1e-12, verbose = TRUE)
pntR1     (t, df, ncp, lower.tail = TRUE, log.p = FALSE,
                           itrmax = 1000, errmax = 1e-12, verbose = TRUE)
pnt3150   (t, df, ncp, lower.tail = TRUE, log.p = FALSE, M = 1000, verbose = TRUE)
pnt3150.1 (t, df, ncp, lower.tail = TRUE, log.p = FALSE, M = 1000, verbose = TRUE)
pntP94    (t, df, ncp, lower.tail = TRUE, log.p = FALSE,
                           itrmax = 1000, errmax = 1e-12, verbose = TRUE)
pntP94.1  (t, df, ncp, lower.tail = TRUE, log.p = FALSE,
                           itrmax = 1000, errmax = 1e-12, verbose = TRUE)
pntLrg    (t, df, ncp, lower.tail = TRUE, log.p = FALSE)
pntJW39   (t, df, ncp, lower.tail = TRUE, log.p = FALSE)
pntJW39.0 (t, df, ncp, lower.tail = TRUE, log.p = FALSE)
pntChShP94  (t, df, ncp, lower.tail = TRUE, log.p = FALSE,
                           itrmax = 1000, errmax = 1e-12, verbose = TRUE)
pntChShP94.1(t, df, ncp, lower.tail = TRUE, log.p = FALSE,
                           itrmax = 1000, errmax = 1e-12, verbose = TRUE)
```

Arguments

t	vector of quantiles (called q in pt(..).
df	degrees of freedom (> 0, maybe non-integer). df = Inf is allowed.
ncp	non-centrality parameter δ; currently except for rt(), only for abs(ncp) <= 37.62. If omitted, use the central t distribution.
log, log.p	logical; if TRUE, probabilities p are given as log(p).
lower.tail	logical; if TRUE (default), probabilities are $P[X \leq x]$, otherwise, $P[X > x]$.
itrmax	...
errmax	...
verbose	...
M	positive integer specifying the number of terms to use in the series.

Details

.: ..

.: ..

pnt3150(), pnt3150.1(): Simple inefficient but hopefully correct version of pntP94..() This is really a direct implementation of formula (31.50), p.532 of Johnson, Kotz and Balakrishnan (1995)

pntP94(), pntP94.1(): New versions of pntR1(), pntR(); using the Posten (1994) algorithm. pntP94() is the Vectorize()d version of pntP94.1().

.: ..

.: ..

Value

a number for pntJKBf1() and .pntJKBch1().

a numeric vector of the same length as the maximum of the lengths of x,df,ncp for pntJKBf() and .pntJKBch().

Author(s)

Martin Maechler

References

Johnson, N.L., Kotz, S. and Balakrishnan, N. (1995) Continuous Univariate Distributions Vol~2, 2nd ed.; Wiley.
Chapter 31, Section *5 Distribution Function*, p.514 ff

See Also

pt, for R's version of non-central t probabilities.

Examples

```
tt <- seq(0, 10, len = 21)
ncp <- seq(0, 6, len = 31)
dt3R   <- outer(tt, ncp, pt, , df = 3)
dt3JKB <- outer(tt, ncp, pntR, df = 3)# currently verbose
stopifnot(all.equal(dt3R, dt3JKB, tolerance = 4e-15))# 64-bit Lnx: 2.78e-16
```

---

ppoisson   *Direct Computation of 'ppois()' Poisson Distribution Probabilities*

Description

Direct computation and errors of ppois Poisson distribution probabilities.

Usage

```
ppoisD(q, lambda, all.from.0 = TRUE)
ppoisErr (lambda, ppFUN = ppoisD, iP = 1e-15,
          xM = qpois(iP, lambda=lambda, lower.tail=FALSE),
          verbose = FALSE)
```

Arguments

q	numeric vector of non-negative integer values, "quantiles" at which to evaluate ppois(q,la) and ppFUN(q,la).
lambda	positive parameter of the Poisson distribution, lambda= $\lambda = E[X] = Var[X]$ where $X \sim Pois(\lambda)$.
all.from.0	logical indicating if q is positive integer, and the probabilities should computed for all quantile values of 0:q.
ppFUN	alternative ppois evaluation, by default the **d**irect summation of dpois(k,lambda).
iP	small number, iP << 1, used to construct the abscissa values x at which to evaluate and compare ppois() and ppFUN(), see xM:
xM	(specified instead of iP:) the maximal x-value to be used, i.e., the values used will be x <-0:iM. The default, qpois(1-iP,lambda = lambda) is the upper tail iP-quantile of Poi(lambda).
verbose	logical indicating if extra information should be printed.

Value

ppoisD() contains the poisson probabilities along q, i.e., is a numeric vector of length length(q).

re <-ppoisErr() returns the relative "error" of ppois(x0,lambda) where ppFUN(x0,lambda) is assumed to be the truth and x0 the "worst case", i.e., the value (among x) with the largest such difference.

Additionally, attr(re,"x0") contains that value x0.

Author(s)

Martin Maechler, March 2004; 2019 ff

See Also

ppois

Examples

```
(lams <- outer(c(1,2,5), 10^(0:3)))# 10^4 is already slow!
system.time(e1 <- sapply(lams, ppoisErr))
e1 / .Machine$double.eps

## Try another 'ppFUN' :--------------------------------
## this relies on the fact that it's *only* used on an 'x' of the form  0:M :
ppD0 <- function(x, lambda, all.from.0=TRUE)
            cumsum(dpois(if(all.from.0) 0:x else x, lambda=lambda))
## and test it:
p0 <- ppD0 (  1000, lambda=10)
p1 <- ppois(0:1000, lambda=10)
stopifnot(all.equal(p0,p1, tol=8*.Machine$double.eps))

system.time(p0.slow <- ppoisD(0:1000, lambda=10, all.from.0=FALSE))# not very slow, here
p0.1 <- ppoisD(1000, lambda=10)
if(requireNamespace("Rmpfr")) {
 ppoisMpfr <- function(x, lambda) cumsum(Rmpfr::dpois(x, lambda=lambda))
 p0.best <- ppoisMpfr(0:1000, lambda = Rmpfr::mpfr(10, precBits = 256))
 AllEq. <- Rmpfr::all.equal
```

```
 AllEq <- function(target, current, ...)
    AllEq.(target, current, ...,
           formatFUN = function(x, ...) Rmpfr::format(x, digits = 9))
 print(AllEq(p0.best, p0,      tol = 0)) # 2.06e-18
 print(AllEq(p0.best, p0.slow, tol = 0)) # the "worst" (4.44e-17)
 print(AllEq(p0.best, p0.1,    tol = 0)) # 1.08e-18
}

## Now (with 'all.from.0 = TRUE',  it is fast too):
p15    <- ppoisErr(2^13)
p15.0. <- ppoisErr(2^13, ppFUN = ppD0)
c(p15, p15.0.) / .Machine$double.eps # on Lnx 64b, see (-10  2.5), then (-2 -2)

## lapply(), so you see "x0" values :
str(e0. <- lapply(lams, ppoisErr, ppFUN = ppD0))

## The first version [called 'err.lambd0()' for years] used simple  cumsum(dpois(..))
## NOTE: It is *stil* much faster, as it relies on special  x == 0:M  relation
## Author: Martin Maechler, Date:  1 Mar 2004, 17:40
##
e0 <- sapply(lams, function(lamb) ppoisErr(lamb, ppFUN = ppD0))
all.equal(e1, e0) # typically TRUE,  though small "random" differences:
cbind(e1, e0) * 2^53 # on Lnx 64b, seeing integer values in {-24, .., 33}
```

---

qbetaAppr *Compute (Approximate) Quantiles of the Beta Distribution*

Description

Compute quantiles (inverse distribution values) for the beta distribution, using diverse approximations.

Usage

```
qbetaAppr.1(a, p, q, y = qnormUappr(a))
qbetaAppr.2(a, p, q, lower.tail=TRUE, log.p=FALSE, logbeta = lbeta(p,q))
qbetaAppr.3(a, p, q, lower.tail=TRUE, log.p=FALSE, logbeta = lbeta(p,q))
qbetaAppr.4(a, p, q, y = qnormUappr(a),
            verbose = getOption("verbose"))

qbetaAppr  (a, p, q, y = qnormUappr(a), logbeta= lbeta(p,q),
            verbose = getOption("verbose") && length(a) == 1)

qbeta.R    (alpha, p, q,
            lower.tail = TRUE, log.p = FALSE,
    logbeta = lbeta(p,q),
    low.bnd = 3e-308, up.bnd = 1-2.22e-16,
            method = c("AS109", "Newton-log"),
            tol.outer = 1e-15,
    f.acu = function(a,p,q) max(1e-300, 10^(-13- 2.5/pp^2 - .5/a^2)),
    fpu = .Machine$ double.xmin,
    qnormU.fun = function(u, lu) qnormUappr(p=u, lp=lu)
```

```
          , R.pre.2014 = FALSE
    , verbose = getOption("verbose")
          , non.finite.report = verbose
          )
```

Arguments

`a, alpha`	vector of probabilities (otherwise, e.g., in [qbeta](), called p).
`p, q`	the two shape parameters of the beta distribution; otherwise, e.g., in [qbeta](), called shape1 and shape2.
`y`	an approximation to $\Phi^{-1}(1-\alpha)$ (aka $z_{1-\alpha}$) where $\Phi(x)$ is the standard normal cumulative probability function and $\Phi{-1}(x)$ its inverse, i.e., R's [qnorm](x).
`lower.tail, log.p`	
	logical, see, e.g., [qchisq](); must have length 1.
`logbeta`	must be [lbeta](p,q); mainly an option to pass a value already computed.
`verbose`	logical or integer indicating if and how much "monitoring" information should be produced by the algorithm.
`low.bnd, up.bnd`	
	lower and upper bounds for ...TODO...
`method`	a string specifying the approximation method to be used.
`tol.outer`	the "outer loop" convergence tolerance; the default 1e-15 has been hardwired in R's [qbeta]().
`f.acu`	a [function]() with arguments (a,p,q) ...TODO...
`fpu`	a very small positive number.
`qnormU.fun`	a [function]() with arguments (u,lu) to compute "the same" as [qnormUappr](), the upper standard normal quantile.
`R.pre.2014`	a [logical]() ... TODO ...
`non.finite.report`	
	[logical]() indicating if during the "outer loop" refining iterations, if y becomes non finite and the iterations have to stop, it should be reported (before the current best value is returned).

Value

...

Author(s)

The R Core Team for the C version in R's sources; Martin Maechler for the R port.

See Also

[qbeta]().

Examples

```
qbeta.R(0.6, 2, 3) # 0.4445
qbeta.R(0.6, 2, 3) - qbeta(0.6, 2,3) # almost 0

qbetaRV <- Vectorize(qbeta.R, "alpha") # now can use
curve(qbetaRV(x, 1.5, 2.5))
curve(qbeta (x, 1.5, 2.5), add=TRUE, lwd = 3, col = adjustcolor("red", 1/2))

## an example of disagreement (and doubt, as borderline, close to underflow):
qbeta.R(0.5078, .01, 5) # -> 2.77558e-15    # but
qbeta (0.5078, .01, 5) # -> 1.776357e-15  now gives 4.651188e-31 !!!
qbeta (0.5078, .01, 5, ncp=0) #    also    gives 4.651188e-31
```

---

qchisqAppr   *Compute Approximate Quantiles of the Chi-Squared Distribution*

Description

Compute quantiles (inverse distribution values) for the chi-squared distribution. using Johnson,Kotz,..
............TODO.......

Usage

```
qchisqKG   (p, df, lower.tail = TRUE, log.p = FALSE)
qchisqWH   (p, df, lower.tail = TRUE, log.p = FALSE)
qchisqAppr  (p, df, lower.tail = TRUE, log.p = FALSE, tol = 5e-7)
qchisqAppr.R(p, df, lower.tail = TRUE, log.p = FALSE, tol = 5e-07,
             maxit = 1000, verbose = getOption("verbose"), kind = NULL)
```

Arguments

p	vector of probabilities.
df	degrees of freedom > 0, maybe non-integer; must have length 1.
lower.tail, log.p	
	logical, see, e.g., qchisq(); must have length 1.
tol	non-negative number, the convergence tolerance
maxit	the maximal number of iterations
verbose	logical indicating if the algorithm should produce "monitoring" information.
kind	the *kind* of approximation; if NULL, the default, the approximation chosen depends on the arguments; notably it is chosen separately for each p. Otherwise, it must be a character string. The main approximations are Wilson-Hilferty versions, when the string contains "WH". More specifically, it must be one of the strings

   "chi.small" particularly useful for small chisquared values p;... ...

   "WH" ... ...

   "p1WH" ... ...

   "WHchk" ... ...

   "df.small" particularly useful for small degrees of freedom df... ...

Value

...

Author(s)

Martin Maechler

See Also

[qchisq](). Further, our approximations to the *non-central* chi-squared quantiles, [qnchisqAppr]()

Examples

```
## TODO
```

---

qgammaAppr   *Compute (Approximate) Quantiles of the Gamma Distribution*

Description

Compute approximations to the quantile (i.e., inverse cumulative) function of the Gamma distribution.

Usage

```
qgammaAppr(p, shape, lower.tail = TRUE, log.p = FALSE,
           tol = 5e-07)
qgamma.R  (p, alpha, scale = 1, lower.tail = TRUE, log.p = FALSE,
           EPS1 = 0.01, EPS2 = 5e-07, epsN = 1e-15, maxit = 1000,
           pMin = 1e-100, pMax = (1 - 1e-14),
           verbose = getOption("verbose"))

qgammaApprKG(p, shape, lower.tail = TRUE, log.p = FALSE)


qgammaApprSmallP(p, shape, lower.tail = TRUE, log.p = FALSE)
```

Arguments

p	numeric vector (possibly log tranformed) probabilities.
shape, alpha	shape parameter, non-negative.
scale	scale parameter, non-negative, see [qgamma]().
lower.tail, log.p	
	logical, see, e.g., [qgamma](); must have length 1.
tol	tolerance of maximal approximation error.
EPS1	small positive number. ...
EPS2	small positive number. ...

epsN	small positive number. ...
maxit	maximal number of iterations. ...
pMin, pMax	boundaries for p. ...
verbose	logical indicating if the algorithm should produce "monitoring" information.

Details

qgammaApprSmallP(p,a) should be a good approximation in the following situation when both p and shape $= \alpha =: a$ are small :

If we look at Abramowitz&Stegun $gamma * (a, x) = x^- a * P(a, x)$ and its series $g * (a, x) = 1/gamma(a) * (1/a - 1/(a + 1) * x + ...)$,

then the first order approximation $P(a, x) = x^a * g * (a, x) = x^a/gamma(a + 1)$ and hence its inverse $x = qgamma(p, a) = (p * gamma(a + 1))^{(1/a)}$ should be good as soon as $1/a >> 1/(a + 1) * x$

<==> x « (a+1)/a = (1 + 1/a)

<==> x < eps *(a+1)/a

<==> log(x) < log(eps) + log((a+1)/a) = log(eps) + log((a+1)/a) ~ -36 - log(a) where log(x) ~= log(p * gamma(a+1)) / a = (log(p) + lgamma1p(a))/a

such that the above

<==> (log(p) + lgamma1p(a))/a < log(eps) + log((a+1)/a)

<==> log(p) + lgamma1p(a) < a*(-log(a)+ log(eps) + log1p(a))

<==> log(p) < a*(-log(a)+ log(eps) + log1p(a)) - lgamma1p(a) =: bnd(a)

Note that qgammaApprSmallP() indeed also builds on [lgamma1p]().

.qgammaApprBnd(a) provides this bound $bnd(a)$; it is simply a*(logEps + log1p(a) -log(a)) -lgamma1p(a), where logEps is $\log(\epsilon) = $ log(eps) where eps <-.Machine$double.eps, i.e. typically (always?) logEps$= \log \epsilon = -52 * \log(2) = -36.04365$.

Value

numeric

Author(s)

Martin Maechler

References

Abramowitz, M. and Stegun, I. A. (1972) *Handbook of Mathematical Functions*. New York: Dover. https://en.wikipedia.org/wiki/Abramowitz_and_Stegun provides links to the full text which is in public domain.

See Also

[qgamma]() for R's Gamma distribution functions.

Examples

```
## TODO :  Move some of the curve()s from ../tests/qgamma-ex.R !!
```

---

qnchisqAppr	*Compute Approximate Quantiles of Noncentral Chi-Squared Distribution*

Description

Compute quantiles (inverse distribution values) for the *non-central* chi-squared distribution.

.... using Johnson,Kotz, and other approximations ....

Usage

```
qchisqAppr.0 (p, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
qchisqAppr.1 (p, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
qchisqAppr.2 (p, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
qchisqAppr.3 (p, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
qchisqApprCF1(p, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
qchisqApprCF2(p, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)

qchisqCappr.2 (p, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
qchisqN       (p, df, ncp = 0, qIni = qchisqAppr.0, ...)

qnchisqAbdelAty  (p, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
qnchisqBolKuz    (p, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
qnchisqPatnaik   (p, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
qnchisqPearson   (p, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
qnchisqSankaran_d(p, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
```

Arguments

p	vector of probabilities.
df	degrees of freedom > 0, maybe non-integer.
ncp	non-centrality parameter δ; ....
lower.tail, log.p	
	logical, see, e.g., qchisq().
qIni	a function that computes an approximate noncentral chi-squared quantile as starting value x0 for the Newton algorithm newton().
...	further arguments to newton(), notably eps or maxiter.

Details

Compute (approximate) quantiles, using approximations analogous to those for the probabilities, see pnchisqPearson.

qchisqAppr.0(): ...TODO...

qchisqAppr.1(): ...TODO...

qchisqAppr.2(): ...TODO...

qchisqAppr.3(): ...TODO...

qchisqApprCF1(): ...TODO...

qchisqApprCF2(): ...TODO...

qchisqCappr.2(): ...TODO...

qchisqN(): Uses Newton iterations with [pchisq](#)() and [dchisq](#)() to determine [qchisq](#)(.) values.

qnchisqAbdelAty(): ...TODO...

qnchisqBolKuz(): ...TODO...

qnchisqPatnaik(): ...TODO...

qnchisqPearson(): ...TODO...

qnchisqSankaran_d(): ...TODO...

Value

[numeric](#) vectors of (noncentral) chi-squared quantiles, corresponding to probabilities p.

Author(s)

Martin Maechler, from May 1999; starting from a post to the S-news mailing list by Ranjan Maitra (@ math.umbc.edu) who showed a version of our qchisqAppr.0() thanking Jim Stapleton for providing it.

References

Johnson, N.L., Kotz, S. and Balakrishnan, N. (1995) Continuous Univariate Distributions Vol~2, 2nd ed.; Wiley.
Chapter 29 *Noncentral χ^2-Distributions*; notably Section *8 Approximations*, p.461 ff.

See Also

[qchisq](#).

Examples

```
pp <- c(.001, .005, .01, .05, (1:9)/10, .95, .99, .995, .999)
pkg <- "package:DPQ"
qnchNms <- c(paste0("qchisqAppr.",0:3), paste0("qchisqApprCF",1:2),
             "qchisqN", "qchisqCappr.2", ls(pkg, pattern = "^qnchisq"))
qnchF <- sapply(qnchNms, get, envir = as.environment(pkg))
for(ncp in c(0, 1/8, 1/2)) {
  cat("\n~~~~~~~~~~~~\nncp: ", ncp,"\n======\n")
  print(sapply(qnchF, function(F) Vectorize(F, "p")(pp, df = 3, ncp=ncp)))
}

## Bug: qnchisqSankaran_d() has numeric overflow problems for large df:
qnchisqSankaran_d(pp, df=1e200, ncp = 100)

## One current (2019-08) R bug: Noncentral chi-squared quantiles on *LOG SCALE*

## a)  left/lower tail : -------------------------------------------------------
qs <- 2^seq(0,11, by=1/16)
pqL <- pchisq(qs, df=5, ncp=1, log.p=TRUE)
plot(qs, -pqL, type="l", log="xy") # + expected warning on log(0) -- all fine
qpqL <- qchisq(pqL, df=5, ncp=1, log.p=TRUE) # severe overflow :
qm <- cbind(qs, pqL, qchisq=qpqL
, qchA.0 = qchisqAppr.0 (pqL, df=5, ncp=1, log.p=TRUE)
```

```
, qchA.1 = qchisqAppr.1 (pqL, df=5, ncp=1, log.p=TRUE)
, qchA.2 = qchisqAppr.2 (pqL, df=5, ncp=1, log.p=TRUE)
, qchA.3 = qchisqAppr.3 (pqL, df=5, ncp=1, log.p=TRUE)
, qchACF1= qchisqApprCF1(pqL, df=5, ncp=1, log.p=TRUE)
, qchACF2= qchisqApprCF2(pqL, df=5, ncp=1, log.p=TRUE)
, qchCa.2= qchisqCappr.2(pqL, df=5, ncp=1, log.p=TRUE)
, qnPatnaik   = qnchisqPatnaik   (pqL, df=5, ncp=1, log.p=TRUE)
, qnAbdelAty  = qnchisqAbdelAty  (pqL, df=5, ncp=1, log.p=TRUE)
, qnBolKuz    = qnchisqBolKuz    (pqL, df=5, ncp=1, log.p=TRUE)
, qnPearson   = qnchisqPearson   (pqL, df=5, ncp=1, log.p=TRUE)
, qnSankaran_d= qnchisqSankaran_d(pqL, df=5, ncp=1, log.p=TRUE)
)

round(qm[ qs %in% 2^(0:11) , -2])
#=> Approximations don't overflow but are not good enough

## b)  right/upper tail , larger ncp -----------------------------------------
qS <- 2^seq(-3, 3, by=1/8)
pqLu <- pchisq(qS, df=5, ncp=100, log.p=TRUE, lower.tail=FALSE)
## using "the alternative" (here is currently identical):
identical(pqLu, (pqLu.<- log1p(-pchisq(qS, df=5, ncp=100)))) # here TRUE
plot (qS, -pqLu, type="l", log="xy") # fine
qpqLu <- qchisq(pqLu, df=5, ncp=100, log.p=TRUE, lower.tail=FALSE)
cbind(qS, pqLu, pqLu, qpqLu)# # severe underflow
qchMat <- cbind(qchisq = qpqLu
, qchA.0 = qchisqAppr.0 (pqLu, df=5, ncp=100, log.p=TRUE, lower.tail=FALSE)
, qchA.1 = qchisqAppr.1 (pqLu, df=5, ncp=100, log.p=TRUE, lower.tail=FALSE)
, qchA.2 = qchisqAppr.2 (pqLu, df=5, ncp=100, log.p=TRUE, lower.tail=FALSE)
, qchA.3 = qchisqAppr.3 (pqLu, df=5, ncp=100, log.p=TRUE, lower.tail=FALSE)
, qchACF1= qchisqApprCF1(pqLu, df=5, ncp=100, log.p=TRUE, lower.tail=FALSE)
, qchACF2= qchisqApprCF2(pqLu, df=5, ncp=100, log.p=TRUE, lower.tail=FALSE)
, qchCa.2= qchisqCappr.2(pqLu, df=5, ncp=100, log.p=TRUE, lower.tail=FALSE)
, qnPatnaik   = qnchisqPatnaik   (pqLu, df=5, ncp=100, log.p=TRUE, lower.tail=FALSE)
, qnAbdelAty  = qnchisqAbdelAty  (pqLu, df=5, ncp=100, log.p=TRUE, lower.tail=FALSE)
, qnBolKuz    = qnchisqBolKuz    (pqLu, df=5, ncp=100, log.p=TRUE, lower.tail=FALSE)
, qnPearson   = qnchisqPearson   (pqLu, df=5, ncp=100, log.p=TRUE, lower.tail=FALSE)
, qnSankaran_d= qnchisqSankaran_d(pqLu, df=5, ncp=100, log.p=TRUE, lower.tail=FALSE)
)
cbind(L2err <- sort(sqrt(colSums((qchMat - qS)^2))))
##--> "Sankaran_d", "CF1" and "CF2" are good here

plot (qS, qpqLu, type = "b", log="x", lwd=2)
lines(qS, qS, col="gray", lty=2, lwd=3)
top3 <- names(L2err)[1:3]
use <- c("qchisq", top3)
matlines(qS, qchMat[, use]) # 3 of the approximations are "somewhat ok"
legend("topleft", c(use,"True"), bty="n", col=c(palette()[1:4], "gray"),
                  lty = c(1:4,2), lwd = c(2, 1,1,1, 3))
```

---

qnormAppr   *Approximations to 'qnorm()', i.e., z_α*

Description

Relatively simple approximations to the standard normal (aka "Gaussian") quantiles, i.e., the inverse
of the normal cumulative probability function.

qnormUappr() is a simple approximation to (the **u**pper tail) standard normal quantiles, [qnorm](https://)().

Usage

```
qnormAppr(p)
qnormUappr(p, lp = .DT_Clog(p, lower.tail=lower.tail, log.p=log.p),
           lower.tail = FALSE, log.p = FALSE)
```

Arguments

p	numeric vector of probabilities, possibly transformed, depending on `log.p`. Does not need to be specified, if `lp` is instead.
lp	`log(1 -p*)`, assuming $p*$ is the `lower.tail=TRUE`, `log.p=FALSE` version of p. If passed as argument, it can be much more accurate than when computed from p by default.
lower.tail	logical; if TRUE (*not* the default here!), probabilities are $P[X \leq x]$, otherwise (by default) upper tail probabilities, $P[X > x]$.
log.p	logical; if TRUE, probabilities p are given as $\log(p)$ in argument p.

Details

qnormAppr(p) uses the simple 4 coefficient rational approximation to [qnorm](https://)(p), to be used *only* for $p > 1/2$ in [qbeta](https://)() computations, e.g., [qbeta.R](https://).
The relative error of this approximation is quite *asymmetric*: It is mainly < 0.

qnormUappr(p) uses the same rational approximation directly for the **U**pper tail where it is relatively good, and for the lower tail via "swapping the tails", so it is good there as well.

Value

numeric vector of (approximate) normal quantiles corresponding to probabilities p

Author(s)

Martin Maechler

See Also

[qnorm](https://).

Examples

```
pp <- c(.001, .005, .01, .05, (1:9)/10, .95, .99, .995, .999)
z_p <- qnorm(pp)
(R <- cbind(pp, z_p, qA = qnormAppr(pp), qUA = qnormUappr(pp, lower.tail=TRUE)))
## Errors, absolute and relative:
cbind(pp, (relE <- cbind(
              errA  = z_p - R[,"qA" ],
              errUA = z_p - R[,"qUA"],
              rE.A  = 1 - R[,"qA" ]/z_p,
              rE.UA = 1 - R[,"qUA"]/z_p)))

lp <- -c(1000, 500, 200, 100, 50, 20:10, seq(9.75, 0, by = -1/8))
qnormUappr(lp=lp) # 'p' need not be specified if 'lp' is
```

```
curve(qnorm(x, lower.tail=FALSE), n=1001)
curve(qnormUappr(x), add=TRUE,    n=1001, col = adjustcolor("red", 1/2))

curve(qnorm(x, lower.tail=FALSE) - qnormUappr(x), n=1001)
```

---

qtAppr   *Compute Approximate Quantiles of Non-Central t Distribution*

Description

Compute quantiles (inverse distribution values) for the non-central t distribution. using Johnson,Kotz,.. p.521, formula (31.26 a) (31.26 b) & (31.26 c)

Note that [qt](..,ncp=*) did not exist yet in 1999, when MM implemented qtAppr().

Usage

```
qtAppr(p, df, ncp, lower.tail = TRUE, log.p = FALSE, method = c("a", "b", "c"))
```

Arguments

p	vector of probabilities.
df	degrees of freedom > 0, maybe non-integer.
ncp	non-centrality parameter δ; ....
lower.tail, log.p	
	logical, see, e.g., [qt]().
method	a string specifying the approximation method to be used.

Value

...

Author(s)

Martin Maechler, 6 Feb 1999

See Also

[qt](.)

Examples

```
## TODO
```

---

r_pois   *Compute Relative Size of i-th term of Poisson Distribution Series*

Description

Compute

$$r_\lambda(i) := (\lambda^i/i!)/e_{i-1}(\lambda),$$

where $\lambda =$ lambda, and

$$e_n(x) := 1 + x + x^2/2! + ... + x^n/n!$$

is the n-th partial sum of $\exp(x) = e^x$.

Questions: As function of i

- Can this be put in a simple formula, or at least be well approximated for large λ and/or large i?

- For which i ($:= i_m(\lambda)$) is it maximal?

- When does $r_\lambda(i)$ become smaller than (f+2i-x)/x = a + b*i ?

NB: This is relevant in computations for non-central chi-squared (and similar non-central distribution functions) defined as weighted sum with "Poisson weights".

Usage

```
r_pois(i, lambda)
r_pois_expr  # the R expression() for the asymptotic branch of r_pois()

plRpois(lambda, iset = 1:(2*lambda), do.main = TRUE,
        log = 'xy', type = "o", cex = 0.4, col = c("red","blue"),
        do.eaxis = TRUE, sub10 = "10")
```

Arguments

i	integer ..
lambda	non-negative number ...
iset	.....
do.main	[logical](#) specifying if a main [title](#) should be drawn via (main = r_pois_expr).
type	type of (line) plot, see [lines](#).
log	string specifying if (and where) logarithmic scales should be used, see [plot.default](#)().
cex	character expansion factor.
col	colors for the two curves.
do.eaxis	[logical](#) specifying if [eaxis](#)() (package **sfsmisc**) should be used.
sub10	argument for [eaxis](#)() (with a different default than the original).

Details

r_pois() is related to our series expansions and approximations for the non-central chisquared; in particular ....

plRpois() simply produces a "nice" plot of r_pois(ii,*) vs ii.

Value

r_pois() returns a numeric vector $r_\lambda(i)$ values.

r_pois_expr() an expression.

Author(s)

Martin Maechler, 20 Jan 2004

See Also

dpois().

Examples

```
plRpois(12)
plRpois(120)
```

Index